

Language Model Behavior: A Comprehensive Survey

Tyler A. Chang
UC San Diego
tachang@ucsd.edu

Benjamin K. Bergen
UC San Diego
bkbergen@ucsd.edu

Transformer language models have received widespread public attention, yet their generated text is often surprising even to NLP researchers. In this survey, we discuss over 250 recent studies of English language model behavior before task-specific fine-tuning. Language models possess basic capabilities in syntax, semantics, pragmatics, world knowledge, and reasoning, but these capabilities are sensitive to specific inputs and surface features. Despite dramatic increases in generated text quality as models scale to hundreds of billions of parameters, the models are still prone to unfactual responses, commonsense errors, memorized text, and social biases. Many of these weaknesses can be framed as over-generalizations or under-generalizations of learned patterns in text. We synthesize recent results to highlight what is currently known about what large language models can and cannot do.

Contents

1	Introduction	1
2	Transformer Language Models	3
3	Syntax	7
4	Semantics and Pragmatics	10
5	Commonsense and World Knowledge	14
6	Logical and Numerical Reasoning	18
7	Memorized vs. Novel Text	19
8	Bias, Privacy, and Toxicity	21
9	Misinformation, Personality, and Politics	25
10	Discussion	27
11	Conclusion	31

1. Introduction

Transformer language models have revolutionized the field of natural language processing (NLP) since their introduction in 2018 (Radford et al. 2018; Devlin et al. 2019). Recent research and public attention has demonstrated that large language models (e.g. GPT-3/4, PaLM, and OPT; Brown et al. 2020; Chowdhery et al. 2022; Zhang et al. 2022b; OpenAI 2023a) can achieve remarkable performance both on standard NLP benchmarks and on open-ended natural language generation tasks from the general public (Wang et al. 2019; Johnson 2022). Already, language models are used in industry for applications ranging from web search and chatbots to medical and financial document analysis (Nayak 2019; Broyde and Palmer 2021; Thewsey 2021; Lee 2023). Due to their widespread applicability, language models have been called “foundation models” for NLP (Bommasani et al. 2021).

Language models are trained to predict masked (i.e. hidden) or upcoming words from context, usually text. The models can then be fine-tuned for specific downstream tasks (e.g. text classification; [Devlin et al. 2019](#)), or they can be used directly for any text prediction task. As language model capabilities have expanded in recent years, they have increasingly been used in the text generation scenario with minimal or no fine-tuning ([Brown et al. 2020](#)). This approach requires no task-specific data or further training infrastructure, thus expanding the range of possibilities and audience for language model applications. In particular, the release of public APIs and interfaces such as GPT-3 and ChatGPT ([Brown et al. 2020](#); [OpenAI 2022](#)) have enabled widespread public experimentation on the text generation capabilities of language models.

Yet, text generated by language models is often surprising even to NLP researchers. Previous studies have investigated both the outputs and internal mechanisms of language models, originally focusing on masked (i.e. fill-in-the-blank) “BERT” models and establishing the field of “BERTology” (see [Rogers, Kovaleva, and Rumshisky 2020](#) for a survey). In the years since the last BERTology survey in 2020, and in tandem with the rise of large autoregressive models such as GPT-3 (i.e. predicting upcoming words instead of masked words), language model analysis has shifted focus to these large autoregressive models. Because these models are often used without fine-tuning for open-ended text generation, there have been an increasing number of behavioral studies evaluating the output text probabilities of language models.

Despite this flurry of research, language model text generation behavior remains unpredictable. Although model performance on broad benchmark datasets is relatively consistent for a given model size and architecture, responses to specific inputs and examples are not. This feature makes large language models tempting but unreliable to use in many practical applications ([Ganguli et al. 2022a](#)). Furthermore, the rapid pace of NLP research and the quantity of individual studies make any progress in understanding model behavior difficult to track. As language models become more widespread and researchers from other fields invest interest in language models, it is increasingly important that our existing understanding of model behavior be made clear and accessible.

In this survey, we discuss over 250 recent studies of English language model behavior, covering syntax, semantics, pragmatics, world knowledge, reasoning, memorization, and bias.¹ Language models generate fluent and coherent text, but their predictions are highly dependent on input context. Slight changes in input word choice and phrasing can lead to unfactual, offensive, or plagiarized text. Understanding these behaviors has broad implications for informed applications in industry ([Weidinger et al. 2021](#)) and general questions about meaning and “understanding” in artificial agents ([Bender and Koller 2020](#); [Mitchell and Krakauer 2022](#); [Shardlow and Przybyla 2022](#)).

To the extent possible, we avoid taking a stance on whether language models truly “understand” language; instead, we hope to provide a thorough review of the empirical evidence for what behaviors language models exhibit in controlled settings. We discuss a wide range of model capabilities and weaknesses (Sections 3 through 9), and we synthesize results framed from the perspectives of model scale (Section 10.1) and text pattern generalization (Section 10.2). In this way, we hope to combat anecdote-driven language model “hype” with informed hype grounded in what language models actually can and cannot do ([Bowman 2022](#)), while also highlighting potential future directions of research in language model behavioral analysis.

¹ The process for identifying papers and studies for this survey is described in Appendix A.

1.1 Scope

We consider studies of masked and autoregressive English Transformer language models not fine-tuned for any specific downstream tasks. We exclude a wealth of research on fine-tuned model behavior (e.g. models tuned for natural language inference, a text classification task). During the fine-tuning process, language models are prone to overfitting to spurious correlations between text features and labels in the fine-tuning dataset (McCoy, Pavlick, and Linzen 2019; Kavumba et al. 2020; Wang et al. 2022b; Du et al. 2022a; Kavumba, Takahashi, and Oda 2022), and they can even “forget” syntactic and semantic information learned during the original pre-training process (Miaschi et al. 2020; Mosbach et al. 2020). Thus, fine-tuned language models are not necessarily reflective of the linguistic abilities of language models in general. Moreover, as noted in the Introduction, language models are increasingly used without fine-tuning on any individual task.

We also leave studies of non-English and multilingual language models to future surveys that can better focus on the many nuances of cross-lingual comparisons. We acknowledge that over-focusing on high-resource languages (e.g. English) is a recurring problem in NLP research (Joshi et al. 2020), and we hope that this survey provides a foundation to expand to less well-studied languages for which language models often perform poorly (Wu and Dredze 2020; Choudhury and Deshpande 2021). Future surveys might also study the behavior of language model variants such as vision-language models (Du et al. 2022b), code models (Chen et al. 2021), speech models (Lakhotia et al. 2021; Radford et al. 2022), knowledge-augmented models (Zhang et al. 2019), sparsely-activated models (Fedus, Zoph, and Shazeer 2022), or compressed models (Sanh et al. 2019; Zafrir et al. 2019). In the current survey, we consider non-augmented “out-of-the-box” Transformer language models, as used in the majority of NLP research.

Finally, we limit our survey to behavioral studies of language models. These studies treat the models as black box functions that take input text and return probability distributions over output text. Often inspired by work in psycholinguistics, these studies evaluate language model responses to controlled inputs (e.g. Ettinger 2020), to make inferences about how the models process and generate text. As we note in Discussion Section 10.3, other studies analyze language models at the mechanistic level, studying internal representations, individual neurons, and attention heads (Geva et al. 2021; Meng et al. 2022; Olsson et al. 2022). We focus on behavioral studies in this survey, but establishing ties between mechanistic and behavioral analyses of language models is an exciting direction of emerging research.

2. Transformer Language Models

In this section, we provide a brief introduction to Transformer language models, which we generally refer to as language models. Transformer language models use a deep neural network architecture called a Transformer (Vaswani et al. 2017; Section 2.1), and they are trained to predict either masked words (i.e. fill-in-the-blank) or upcoming words in text (Section 2.2). Throughout this survey, we refer to these two types of models as masked and autoregressive models respectively.² Some studies refer to them

² Along with differentiating results for masked vs. autoregressive models, we mention when studies use a GPT-3 model (autoregressive) that may or may not have been instruction-tuned (Section 2.2). For example, `text-davinci-001` and `text-davinci-002` are instruction-tuned, but `davinci` is not (OpenAI 2023b). Still, even the instruction-tuning stage uses only the language modeling

as bidirectional and unidirectional models. Language models are most often applied to downstream tasks using either fine-tuning (or prompt-tuning), zero-shot prompting, or few-shot prompting (Section 2.3).

2.1 Architectures

The basic Transformer language model architecture has remained largely unchanged since 2018 (Radford et al. 2018; Devlin et al. 2019). First, an input text string is converted into a sequence of tokens. Tokens correspond roughly to words, although some words are composed of multiple subword tokens due to limited vocabulary size. For example, the string “This is preposterous!” might be tokenized into `[_this, _is, _prepo, ster, ous, !]`. Common tokenization techniques include byte pair encoding (BPE; Sennrich, Haddow, and Birch 2016) and unigram language modeling (Kudo 2018), but we refer to these other papers for detailed descriptions of tokenization techniques. Model vocabularies generally range from 30K to 250K possible tokens (Radford et al. 2019; Cohen et al. 2022; Chowdhery et al. 2022).

After tokenization, each token is mapped to a fixed vector “embedding”; the embedding for each token is learned during the pre-training process. The embeddings are passed through a stack of Transformer layers (Vaswani et al. 2017; usually 10-100 layers), each consisting of a self-attention network, layer normalizations, and feedforward networks. The primary innovation of Transformer layers is the self-attention network, which “mixes” the sequence of token embeddings using projections into a “query”, “key”, and “value” vector for each token. This mixing of token embeddings results in a “contextualized” representation for each token, essentially a vector representation that incorporates the context of the input sequence. Finally, after the stack of Transformer layers, each output token representation is projected into a distribution over the same token vocabulary used in the input. In other words, the overall architecture maps each input token to a probability distribution over output tokens (e.g. upcoming tokens). Language models usually have between 100M and 500B total parameters, with autoregressive models usually much larger than masked models (Devlin et al. 2019; Brown et al. 2020; Lieber et al. 2021; Smith et al. 2022b; Chowdhery et al. 2022).

The Transformer architecture does not naturally encode any information about each token’s position in an input sequence; intuitively, it is useful to encode this information for features such as word order. Thus, Transformer language models use a variety of position encoding techniques (Wang et al. 2021a; Dufter, Schmitt, and Schütze 2022), such as adding absolute position embeddings to the input token embeddings (i.e. an embedding for each position i ; Vaswani et al. 2017; Radford et al. 2018; Devlin et al. 2019; Radford et al. 2019; Brown et al. 2020; Zhang et al. 2022b), relative position embeddings or biases (i.e. encoding relative position distances between tokens; Shaw, Uszkoreit, and Vaswani 2018; Dai et al. 2019; Raffel et al. 2020; Chang et al. 2021; Rae et al. 2021; Cohen et al. 2022), or rotary position embeddings (an efficient approach to relative position biases; Su et al. 2021; Chowdhery et al. 2022). With relative rather than absolute position methods, language models can better extrapolate to longer sequences than observed during pre-training (Press, Smith, and Lewis 2022). Language

objective. We specifically note if any study uses a model tuned with reinforcement learning (Section 2.2), e.g. `text-davinci-003`. When we refer to masked and autoregressive language models generally, we refer to models that are not fine-tuned.

models are usually pre-trained with input sequence lengths of around 500 to 2000 tokens.

2.2 Training

Language modeling refers to predicting tokens from context, usually text. Masked and autoregressive language models are pre-trained to predict masked (i.e. hidden) and upcoming tokens respectively. Recall from the previous section that the Transformer architecture predicts an output token distribution for each input token.

The [MASK] walked. (1)

The ____
The dog ____
The dog walked ____ (2)

In masked language models (Example 1), randomly selected tokens are replaced with [MASK] tokens; for each input [MASK] token, the model produces a probability distribution over the token that was masked (i.e. fill-in-the-blank). In autoregressive models (Example 2), no tokens are replaced; for each input token, the model produces a probability distribution over the next token (i.e. predicting each next token).

Language models are pre-trained using gradient descent, observing many examples as in Examples 1 and 2. Text corpora for pre-training usually range from approximately 5B to 1.5T tokens (roughly 15GB to 5TB of raw text; [Devlin et al. 2019](#); [Liu et al. 2019](#); [Brown et al. 2020](#); [Rae et al. 2021](#); [Hoffmann et al. 2022](#)). For compute-optimal pre-training in autoregressive language models, as the number of model parameters increases, the number of pre-training tokens should increase roughly proportionally ([Kaplan et al. 2020](#); [Hoffmann et al. 2022](#)). During pre-training, examples are fed into the models with anywhere from 100K to 4M tokens per optimization step (i.e. batch size), usually with larger batch sizes in larger models ([Devlin et al. 2019](#); [Brown et al. 2020](#); [Hoffmann et al. 2022](#); [Chowdhery et al. 2022](#); [Zhang et al. 2022b](#)). Models are usually pre-trained for 100K to 1M steps ([Radford et al. 2018](#); [Devlin et al. 2019](#); [Zhang et al. 2022b](#)); when possible, examples are not repeated during pre-training ([Hoffmann et al. 2022](#); [Chowdhery et al. 2022](#)). Due to high computational costs, relatively few language models are pre-trained from scratch as described here, and they are usually trained in industry labs. In practice, most NLP researchers build applications upon existing pre-trained language models, using the approaches described in Section 2.3.

This survey considers pre-trained language models as described above. Recent language models often contain further non-task-specific fine-tuning stages (particularly autoregressive models; [Cohen et al. 2022](#); [Ouyang et al. 2022](#)). For example, autoregressive models are sometimes fine-tuned using the language modeling objective on curated human-written examples that demonstrate desirable text outputs ([Ouyang et al. 2022](#)) or examples of outputs that correctly follow input instructions ([Wei et al. 2022a](#); [Iyer et al. 2022](#)). These approaches are referred to as supervised fine-tuning (SFT) or instruction tuning. Some more recent models are also tuned using reinforcement learning, with predicted human preferences for different responses used as a reward (reinforcement learning from human feedback, or RLHF; [Ouyang et al. 2022](#); [OpenAI 2023a](#)). Throughout this survey, we consider non-fine-tuned language models unless otherwise

specified.³ Non-fine-tuned language models still serve as the foundation for more recent language models.

2.3 Downstream tasks and text generation

Language models are used for a wide range of downstream tasks, including but not limited to custom chatbots, question answering, sentiment classification, offensive text detection, and textual similarity quantification (Devlin et al. 2019; Zhang et al. 2020; Zhao, Zhang, and Hopfgartner 2021; Zong and Krishnamachari 2022). Traditionally, given example inputs and outputs for a task, language models are fine-tuned by adjusting all or some model parameters using gradient descent (Radford et al. 2018; Devlin et al. 2019; Lester, Al-Rfou, and Constant 2021; Chowdhery et al. 2022). As autoregressive models have risen in popularity, tasks are increasingly formulated as prompted text generation tasks (Wei et al. 2022a):

Premise: Fun for adults and children.
 Hypothesis: Fun for only children.
 Does the premise entail the hypothesis? (3)

————— (Williams, Nangia, and Bowman 2018)

The input text is referred to as the prompt or context. Autoregressive language models can perform many tasks similar to Example 3 without fine-tuning on that specific task (i.e. zero-shot learning, e.g. by instruction-tuning on other tasks; Wei et al. 2022a). If example inputs and outputs (e.g. 1-100 examples) are included in the prompt, then language models can perform well without any fine-tuning at all (Brown et al. 2020; Chowdhery et al. 2022; Zhang et al. 2022b); providing examples in context without any parameter updates is commonly known as few-shot prompting or in-context learning.

In cases such as Example 3, autoregressive language models can compute the probability for any desired output text by iteratively multiplying the probability for each next token. When the models are used for open-ended text generation (i.e. the models must select each next token), common approaches are to (1) iteratively select the most probable next token (greedy sampling), (2) iteratively sample the next token from the output probability distribution with some temperature parameter τ (temperature sampling), (3) sample from the top k token predictions (top- k sampling), or (4) sample from the top tokens that sum to some probability p (nucleus sampling; Holtzman et al. 2020). In all of these cases, multiple candidate sequences of tokens can be generated and then ranked according to their overall sequence probability (i.e. beam search; Freitag and Al-Onaizan 2017), but beam search is often not used in practice due to its high computational cost. Of the studies discussed in this survey, the majority use greedy, temperature, top- k , or nucleus sampling for open-ended text generation. In the next sections, we discuss recent studies evaluating language model generated text and output text probabilities from a wide range of perspectives.

³ Mentions of GPT-3 specifically may be instruction-tuned, but not tuned with reinforcement learning. See footnote in Section 2.

3. Syntax

We begin with studies that evaluate language model predictions from a syntactic perspective. In the vast majority of cases, language models are more likely to predict grammatical tokens than ungrammatical tokens, adhering to a wide variety of syntactic rules (Section 3.1). In subject-verb agreement, the models’ performance degrades in more complex or infrequent examples (Section 3.2), and language model predictions are possibly over-sensitive to token position information (i.e. word order; Section 3.4), but syntactic abilities overall are learned fairly robustly early in pre-training (Section 3.3).

3.1 Language models generally produce grammatical text.

Systematic syntactic evaluations of autoregressive language models are conducted in Warstadt et al. (2020), Hu et al. (2020), and Gauthier et al. (2020), comparing model probabilities for minimal pair examples that differ in grammaticality due to just one token (e.g. “the boy [*eat/eats]”).⁴ Similar assessments are run for masked language models in Park, Park, and Song (2021). Both autoregressive and masked language models consistently assign higher probabilities to grammatical tokens, and they make predictions consistent with hierarchical syntactic structure, where clauses can be nested within one another. Such structures are commonly observed in human language (Carnie 2002), creating token relationships that are not solely dependent on linear word order.

The girl who had three dogs [*play/plays] accordion. (4)

In Example 4, replacing “girl” with “girls” would require the verb to change to “play”. In other words, the verb “plays” agrees in number with the noun “girl” despite the appearance of the nested clause “who had three dogs” including the distractor noun “dogs” closer to the verb. In these long-distance subject-verb agreement examples, language models generally assign higher probabilities to grammatical options, but their performance varies depending on the specific nouns, verbs, and distractors involved (Section 3.2).

Outside of agreement, language models recognize licensing, when the grammaticality of a token depends on an upstream “licensor” token, usually equal or higher in the hierarchical syntactic structure.

I know what the lion devoured [*the gazelle/_] yesterday.
I know that the lion devoured [the gazelle/ * _] yesterday. (5)

(Wilcox, Futrell, and Levy 2022)

In Example 5, the word “what” licenses the omitted direct object “gazelle” for the verb “devoured”; the word “that” does not license such an omission. This omission licensing is known as a filler-gap dependency, and Wilcox, Futrell, and Levy (2022) find that autoregressive language models respect filler-gap rules. Similarly, masked language models assign higher probabilities to licensed tokens in reflexive licensing (reflexives such as “himself” require a properly situated previous noun phrase; Hu, Chen, and Levy 2020) and in negative polarity items (NPIs such as “any” require a previous negative

⁴ An asterisk before a phrase indicates ungrammaticality, as in Carnie (2002).

word such as “not”; Warstadt et al. 2019). However, autoregressive model predictions for reflexive licensing are less accurate in sentences where the licensed reflexive depends on the specific verb involved (Lee and Schuster 2022).⁵

In general, the grammaticality of language model predictions improves with model size and pre-training corpus size, in both autoregressive and masked models (Warstadt et al. 2020; Pérez-Mayos, Ballesteros, and Wanner 2021). Across model sizes, better overall language modeling performance (e.g. inverse perplexity) is positively correlated with syntactic ability, although this relationship is not clear within any given model size (Hu et al. 2020; Pérez-Mayos, Ballesteros, and Wanner 2021). That said, many syntactic rules may be learned primarily based on memorized examples, dependent on the specific words and structures seen during pre-training (Section 3.2). For example, in cases where people generate syntactically anomalous phrases (e.g. article-noun disagreement between “a” and “days” in “a cold five days”), GPT-3 acceptability predictions roughly mirror human judgments (Mahowald 2023).⁶ When prompted with examples, GPT-3 can answer questions directly about a sentence’s syntactic structure (Zhang et al. 2022a). The results in this section demonstrate basic syntactic abilities in language models.

3.2 Language models learn subject-verb agreement, but they are sensitive to intervening clauses and specific words.

Language models’ syntactic abilities are most often evaluated using agreement, when one token’s form depends on a property of another token. For example, subject nouns in English must agree in number with their corresponding verbs (e.g. “the dog eats” vs. “the dogs eat”; see also Example 4). Masked and autoregressive language models are generally good at predicting verb forms for subject-verb agreement (van Schijndel, Mueller, and Linzen 2019), even in nested clauses and with long-distance dependencies as in Example 4 (Goldberg 2019). However, agreement performance degrades as the distance between the subject and verb increases (Bacon and Regier 2019; Ryu and Lewis 2021; Lakretz et al. 2022). In large autoregressive models, this degradation can be reduced significantly if models are provided with even just two initial examples (using few-shot prompting), as human raters usually are (Lampinen 2022).

Subject-verb agreement performance in language models is also dependent on the specific nouns and verbs involved (Yu et al. 2020; Chaves and Richter 2021). Masked and autoregressive models produce over 40% more accurate agreement predictions for verbs that are already probable from context (Newman et al. 2021), and agreement accuracy is worse overall for infrequent verbs (Wei et al. 2021). For infrequent verbs, masked language models are biased towards the more frequent verb form seen during pre-training (e.g. singular vs. plural) (Wei et al. 2021). Error rates exceed 30% for infrequent verbs in nonce (grammatically correct but semantically meaningless) sentences (Wei et al. 2021), with further degradations if there is an intervening clause between the subject and verb as in Example 4 (Lasri, Lenci, and Poibeau 2022a). This subject-verb agreement degradation in nonce sentences with long-distance dependencies has

⁵ Specifically, Lee and Schuster (2022) study subject- and object-control verbs, as in the sentences:

“The artist promised the lawyers to make fun of [himself/*themselves].”

“The artist persuaded the lawyers to make fun of [*himself/themselves].”

⁶ Acceptability predictions in Mahowald (2023) are elicited from GPT-3 using few-shot prompting (Section 2.3).

also been observed in people, although to a lesser degree than in language models (Lasri et al. 2022). Finally, subject-verb agreement performance in masked and autoregressive language models is dependent on the specific subject noun, although these differences in performance do not appear to be driven by noun frequency (Yu et al. 2020). In many ways, language models’ variable performance on subject-verb agreement reflects a larger sensitivity to specific words and input structures (Discussion Section 10.2).

3.3 Language models learn syntactic rules early in pre-training.

The acquisition of syntactic rules is fairly consistent during language model pre-training. Syntactic rules are learned within roughly the first 20% of masked language model pre-training, as measured by the syntactic generalization suites in Section 3.1 (Liu et al. 2021; Zhang et al. 2021b). Small masked language models (8M parameters) pre-trained on only 30M words of transcribed child-directed speech can achieve similar syntactic performance to standard masked models with over 10x more parameters and 1000x more pre-training data (Huebner et al. 2021). Autoregressive and masked models tend to learn similar syntactic generalizations during the pre-training process regardless of random initializations and training data shuffling (Choshen et al. 2022; Misra 2022). Early in pre-training, models are syntactically more similar to bag-of-words, unigram, and n -gram models (Choshen et al. 2022), passing through stages where their predictions mirror unigram then bigram distributions (Chang and Bergen 2022).⁷ Notably, syntactic abilities emerge in Transformer language models despite the fact that Transformers cannot model arbitrarily deep hierarchical structures unless their number of layers or attention heads increases with input length (Hahn 2020), and Transformers have a tendency to generalize linearly rather than hierarchically when trained from scratch on purely syntactic tasks (Petty and Frank 2021).

3.4 Language models can learn word order without explicit position information, but word order is not necessary in many examples.

At first glance, language modeling performance would seem highly dependent on a model’s understanding of word order (i.e. token positions). For example, syntactic information in English is largely determined by token positions (e.g. “the dog saw the cat” vs. “the cat saw the dog”). However, masked language models pre-trained on data with shuffled words can still be fine-tuned for reasonable performance on a variety of downstream tasks (Sinha et al. 2021). This result may be because token position embeddings (Section 2.1) are still learned through common subword token sequences that remain unshuffled. Even when pre-training data is shuffled after tokenization, masked models learn informative position embeddings using correlations between sentence length and token frequencies (Abdou et al. 2022). Similarly, autoregressive language models without any position embeddings are able to encode token position information implicitly by “counting” the previous tokens in the causal (autoregressive) attention

⁷ Bag-of-words models only have access to surrounding tokens without any word order information. Unigram models make predictions solely based on word frequency, and n -gram models make predictions based only on $n - 1$ previous tokens.

mask (Haviv et al. 2022).⁸ Thus, to some degree, the models in these studies are still able to rely on learned token position information.

In contrast, token position information is removed entirely in masked language models when position embeddings are removed. Small masked language models (e.g. 13M parameters) achieve similar language modeling performance when pre-trained with and without position embeddings, particularly if few tokens are masked per sequence (Chang et al. 2021; Lasri, Lenci, and Poibeau 2022b). However, more masking during pre-training improves fine-tuning performance for larger masked models (Wettig et al. 2023); in these larger models, removing token position information entirely might lead to more detrimental effects than in smaller models. While position information (word order) is not necessary for disambiguating semantic meaning in many sentences, there exists a minority of cases where position cues are necessary (Mahowald et al. 2022). Language models can reconstruct text from shuffled inputs, but not with perfect accuracy (Malkin et al. 2021). Thus, high performing models likely need to learn token position information without overfitting to irrelevant position cues. Both masked and autoregressive models with absolute position embeddings (Section 2.1) exhibit such overfitting, making worse language modeling predictions when sequences are shifted by a constant (i.e. shifting all positions by k , maintaining relative positions), a transformation that would ideally have little effect (Sinha et al. 2022b). This overfitting to position cues may also be related to language models’ tendency to generate highly frequent local structures (shorter n -grams based on local positions) rather than long-term coherent text, as described in Section 7.2.

4. Semantics and Pragmatics

On top of syntax, language models display basic semantic abilities, considering how text can be parsed to produce “meaning”. Language models learn word meanings and relationships as reflected in lexical semantics (Section 4.1), they track entities in described situations (Section 4.3), and they recognize basic figurative language (Section 4.4). However, they struggle with negation (Section 4.2) and pragmatics (Section 4.5).

We begin with compositional and formal semantics, where words and phrases combine in systematic ways to produce novel “meanings”, or at least coherent text. There are relatively few behavioral studies of phrase-level compositionality in non-fine-tuned language models (Hupkes et al. 2022), likely because assessments of how models combine phrases to construct meaning are difficult to study behaviorally without a downstream task.

Camila gave a cake in storage to Emma.

→ give(agent=Camila, theme=cake(nmod.in=storage), recipient=Emma) (6)

(Qiu et al. 2022)

When provided with examples (few-shot prompting; see Section 2.3), autoregressive language models can extract compositional semantic parses from sentences as in Example 6, with performance improving with model size (Qiu et al. 2022; Hosseini et al. 2022). However, because the models are explicitly asked for a semantic parse and the

⁸ The causal attention mask in autoregressive language models only allows tokens to “attend” to previous tokens in the input. Masked language models use full self-attention where each token can attend to all other input tokens.

task output is not natural English, it remains unclear whether and how language models construct “meaning” in more natural scenarios.

4.1 Language models learn semantic and compositional properties of individual words, including argument structure, synonyms, and hypernyms.

Researchers have primarily evaluated compositional semantics in language models through the lens of lexical semantics, which studies word meanings and relationships, considering how individual words influence the meaning and semantic structure of a phrase (Geeraerts 2017). At the word meaning level, both masked and autoregressive language models can predict frequent words from their definitions and vice versa, but they struggle with infrequent words (Senel and Schütze 2021). Masked models can predict noun hypernyms (e.g. “robins” are “birds”) using template sentences (e.g. “A robin is a _”; Hanna and Mareček 2021) or by predicting noun replacements (Ravichander et al. 2020), but predictions degrade when the noun is plural or the hypernym pair is infrequent. The hypernym prediction confidence in autoregressive and masked models is correlated with the human-rated typicality of the hyponym within the hypernym category, with larger models showing stronger typicality effects (Misra, Ettinger, and Rayz 2021). When predicting masked nouns more generally, masked language models assign high probabilities to word synonyms and co-hyponyms (e.g. “robin” and “sparrow” are co-hyponyms of “bird”), rather than pairs of hyponyms and hypernyms (Arefyev et al. 2020). These results suggest that language models understand basic word meanings and allowable word substitutions; more grounded knowledge of the objects and entities that words refer to, such as physical properties and facts, are discussed in Section 5.

Lexical semantics also considers how words influence semantic structure within a clause. Autoregressive models are more likely to predict verbs in the correct argument structure (e.g. the correct number and type of arguments in “gave” in Example 6), but with less accuracy than many syntactic tasks (Warstadt et al. 2020).

Sally frightened Mary because she was so terrifying.
Sally feared Mary because she was so terrifying.

(7)

(Davis and van Schijndel 2020)

Specifically, many studies consider implicit causality in verbs. In Example 7, the verb “frightened” biases the next clause to refer to the verb subject “Sally”. The verb “feared” biases the next clause to refer to the verb object “Mary”. After observing an implicit causality verb, autoregressive models with 1.5B parameters are more likely to predict pronoun genders matching the subject vs. object causality bias of the verb (Davis and van Schijndel 2020); however, this effect only sometimes replicates in masked and autoregressive models under 1B parameters (Upadhye, Bergen, and Kehler 2020; Kementchedjheva, Anderson, and Søgaard 2021). Predictions in these smaller autoregressive models match human verb causality biases more closely for frequent verbs (Huynh, Lentz, and van Miltenburg 2022). Outside of implicit causality, masked and autoregressive models predict prepositional vs. double-object dative alternations (e.g. “gave the book to her” vs. “gave her the book”) according to verb-specific biases, with higher correlations with human ratings in larger models (Hawkins et al. 2020). These verb-specific effects in language models demonstrate a basic understanding of how verb properties affect upcoming syntactic and semantic structures.

4.2 Language models struggle with negation, often performing worse as models scale.

One notable example of compositionality is negation, where a word such as “not” inverts the meaning of a phrase. Masked language models often ignore negation when producing completions, such that they are more likely to generate incorrect completions than correct completions to negated primes (e.g. “A robin is not a [bird]”; (Ettinger 2020; Kassner and Schütze 2020). In fact, autoregressive models generate more incorrect completions after “few”-type quantifiers (e.g. “Few robins are [birds]”) as models increase in size (Michaelov and Bergen 2022b). These results may reflect a similarity to human on-line processing (e.g. neural responses and reading times) rather than offline processing and reasoning (Michaelov and Bergen 2022b). Sensitivity to negation can be improved if language models are fine-tuned on more negation sentences, still using the language modeling objective (predicting tokens); masked models are then much less likely to predict any token that was negated in a given context (Gubelmann and Handschuh 2022).

Negation degrades language model performance in tasks involving more explicit reasoning as well (e.g. reasoning abilities in Section 6). When autoregressive models are presented with negated task prompts (e.g. “Please produce a possible *incorrect* answer to the question”), they perform worse as they increase in size (Jang, Ye, and Seo 2022). Performance is often over 50% worse on negated prompts compared to the original prompts. These weaknesses may not be reflected in many NLP benchmarks due to underrepresentation of negation relative to naturally occurring corpora, and the fact that negation is not relevant for many examples (Hossain, Chinnappa, and Blanco 2022); fine-tuned language models perform much worse on datasets that explicitly focus on negation (Hossain et al. 2020; Geiger, Richardson, and Potts 2020; Tejada, Scholtes, and Spanakis 2021; Truong et al. 2022).

4.3 Language models construct coherent but brittle situation models.

Similar to situation models proposed in human language comprehension (Zwaan 2016), language models are able to track entities such as objects and characters throughout a passage. Autoregressive models are able to recognize whether a phrase introduces a new entity (e.g. the “cake” in “I saw Michael bake a cake” vs. “I doubt Michael baked a cake”), with better accuracy in larger models (Schuster and Linzen 2022). However, when multiple nouns are present, the models sometimes refer to un-introduced entities (e.g. “I doubt Michael baked a cake. It’s in the oven.”; Schuster and Linzen 2022). Masked language models are able to predict the antecedents of bridging anaphora, when an entity (e.g. “the window”) has an implied relation to a previously-mentioned entity (e.g. “the house”) (Pandit and Hou 2021).

When prompted with a passage, GPT-3 can answer questions about entity states and event likelihoods, but only marginally better than chance (Zhang et al. 2023). GPT-3 performs better when answers are stated explicitly in the passage, but its answers are sensitive to the phrasing of the question (Summers-Stay, Bonial, and Voss 2021). GPT-3 also has poor accuracy for questions that involve mathematical reasoning, temporal ordering of events, or logical negation (Summers-Stay, Bonial, and Voss 2021; see also Section 4.2 for negation and Section 6.2 for numerical reasoning). Of course, the studies above consider entities and entity states that are described relatively unambiguously in the text, and language models already exhibit somewhat unreliable performance; in

later sections, we discuss commonsense inferences about the implied mental states of characters (Section 4.5) and implied relationships between events (Section 5.3).

4.4 Language models recognize basic analogies, metaphors, and figurative language.

Contradicting the rules of compositional semantics (Section 4), some phrases have meanings that cannot be constructed directly from their constituent words. Common examples of noncompositional expressions include analogies, metaphors, and idioms; these expressions must be interpreted nonliterally (i.e. figuratively or metaphorically). Masked language models assign higher probabilities to literal sentences, then conventional (i.e. common) metaphors, then novel metaphors, then nonsense (Pedinotti et al. 2021a; Griciūtė, Tanti, and Donatelli 2022). When prompting autoregressive models directly to identify metaphorical language, the models exhibit a sharp increase in performance around 100B parameters (Comşa, Eisenschlos, and Narayanan 2022). From these results, it appears that language models recognize metaphorical language to some degree as they increase in size.

Furthermore, masked and autoregressive models can predict the correct interpretations of similes (figurative comparisons using “like” or “as”), with improvements based on model size, but consistently worse than people (Liu et al. 2022a; He et al. 2022a). The models can complete analogies (e.g. “X is to Y as Z is to _”) reasonably well (Ushio et al. 2021), but they perform significantly worse for more abstract and unconventional analogies (Czinczoll et al. 2022). GPT-3 can generate analogies of comparable quality to people when given open-ended prompts (e.g. “What is analogous to X?”), although quality varies by prompt template (Bhavya, Xiong, and Zhai 2022).

Finally, noncompositional expressions include constructions, linguistic templates whose meanings are not necessarily built up from their constituent words. For example, the comparative correlative construction (e.g. “the better your syntax, the better your semantics”) has a well-understood meaning in English despite its apparent ungrammaticality (e.g. no inflected verb). Masked language models struggle to recognize the comparative correlative, making inferences about the implied descriptions at chance level after accounting for adjective frequencies (Weissweiler et al. 2022). However, research on a wider range of constructions is necessary to determine which constructions language models struggle with more generally.

4.5 Language models can infer the mental states of characters in text, but they struggle with implied meaning and pragmatics.

The previous sections focused on linguistic structure and meaning somewhat independent of context. In conversation, many utterances have implied meanings that depend on context and the intentions of the speaker; these meanings are the focus of pragmatics. According to Grice’s maxims of conversation (quantity, quality, relation, and manner), utterances should be appropriately informative, true, relevant, and clear (Grice 1975). Comprehending and producing pragmatically sound utterances likely requires some sensitivity to others’ mental states (Frank and Goodman 2012; Monroe and Potts 2015; Sikos et al. 2021). Indeed, when asked directly, GPT-3 can infer the knowledge and desires of characters in text (Summers-Stay, Bonial, and Voss 2021; Sap et al. 2022), and it can explain why characters perform actions in everyday situations based on commonsense reasoning (Lal et al. 2022). It can even answer questions about characters’ deceit, indirect requests, irony, implied meaning, and humor, but this ability is not observed in smaller autoregressive models (e.g. 100M parameters) (Hu et al. 2022). When using

a fill-in-the-blank word prediction task to infer knowledge states of characters (e.g. whether they know the location of an object), GPT-3 performs well above chance but worse than people (Trott et al. 2022). Masked language models can predict “go” vs. “come” in narratives with accuracy similar to people, recognizing the implied spatial perspective of the narrative (Masis and Anderson 2021).

However, sensitivity to perspectives and mental states does not translate directly into pragmatic understanding in language models. Autoregressive models are more likely to repeat an entity (e.g. “the cup”) than use a pronoun (e.g. “it”) in many cases where a pronoun would be more natural, thus producing potentially over-informative text (Beyer, Loáiciga, and Schlangen 2021). When explicitly interpreting pragmatically implied meanings (implicatures, e.g. “A asked X, and B responded Y, which means [yes/no]”), both masked and autoregressive models perform only slightly above chance and much worse than people, with no substantial improvements using larger models (Ruis et al. 2022). GPT-3 is unable to predict plausible presuppositions (e.g. “Grant stopped eating meat” implies “Grant once ate meat”) or scalar implicatures (e.g. “some brothers” implies “not all brothers”) any better than chance (Cong 2022). This is in line with studies showing that fine-tuned language models rely on surface cues such as specific function words when they appear to recognize presuppositions (Kabbara and Cheung 2022). That said, both masked and autoregressive models prefer conversationally-relevant content over less relevant content, preferring to output text related to main clause content over embedded clause content (Kim, Yu, and Ettinger 2022). In other words, language models exhibit reasonable sensitivity to relevance and mental states, but their pragmatic abilities struggle overall.

5. Commonsense and World Knowledge

Beyond their ability to interpret and produce fluent text, language models exhibit basic world knowledge, including commonsense reasoning and facts. They learn encyclopedic facts and commonsense properties of objects (Section 5.1), albeit unreliably (Section 5.2), and they have a limited ability to infer typical relationships between actions and events (Section 5.3).

5.1 Language models learn facts and commonsense properties of objects, particularly as models scale, but they are less sensitive than people to physical properties.

Masked and autoregressive language models assign higher probabilities to facts than to alternatives when expressed as sentences (e.g. the knowledge triple in Example 8) (Davison, Feldman, and Rush 2019; Petroni et al. 2019).

Knowledge triple: (Dante, born-in, Florence)
 Natural language template: X was born in Y.
 → Fill-in-the-blank sentence: Dante was born in ____.

(8)

(Petroni et al. 2019)

Language models can complete these sentences for a wide variety of facts, covering countries and locations, popular products, historical figures, and even genres of books, movies, and music (Petroni et al. 2019; Penha and Hauff 2020). This ability improves if researchers use better fill-in-the-blank template sentences, such as naturally-occurring

templates from Wikipedia (Jiang et al. 2020b), or if templates are paired with some relevant preceding context (Adolphs, Dhuliawala, and Hofmann 2021).

However, autoregressive models perform worse when considering larger sets of facts in open-ended factual question-answering (Kalo and Fichtel 2022). Masked and autoregressive models perform poorly when predicting numeric literals (e.g. years; Kalo and Fichtel 2022) and numerical commonsense (e.g. “A bird has _ legs”; Lin et al. 2020) (see Section 6.2 for more general numerical reasoning). The models also struggle to make fine-grained property distinctions between related concepts and hypernyms (e.g. properties of “robins” vs. “birds” in general), although accuracy improves with model size (Peng et al. 2022; Misra, Rayz, and Ettinger 2023). As model size increases, autoregressive models are also more likely to correctly use their background factual knowledge to answer questions; accuracy on relevant facts is more predictive of a correct response to a target question in larger models (Sahu et al. 2022). On top of generally higher accuracy (Kalo and Fichtel 2022), larger models (e.g. 50B parameters) are able to assess whether their own answers to factual questions are correct or incorrect, with this self-reflection ability increasing with model size (Kadavath et al. 2022).

To some degree, language models are also able to predict physical properties of objects, such as colors and sizes, using templates similar to Example 8. Perhaps unsurprisingly, model predictions are generally less sensitive than human responses to real world physical properties. For example, masked models can predict typical vs. atypical properties when prompted using quantifiers (e.g. “All X are _” vs. “Some X are _”; Apidianaki and Garí Soler 2021). However, their property predictions are only loosely correlated with human responses, and when predicting a target object from its properties, the models rely on encyclopedic facts over visual and perceptual properties (Weir, Poliak, and Durme 2020). Both masked and autoregressive models can predict typical color distributions of objects, but their predictions correlate more with corpus *n*-grams (e.g. “red ball”) than with human judgments (Paik et al. 2021), particularly for smaller models (Liu et al. 2022b). Similarly, autoregressive models assign higher probabilities to correct physical comparisons (e.g. “A bear is bigger than a cat”) than to incorrect comparisons, with better performance in larger models (Shi and Wolff 2021; De Bruyn et al. 2022). Finally, masked models can predict the typical use for an object better than chance (Jiang and Riloff 2021), and GPT-3 predicts atypical but physically plausible (i.e. “afforded”) uses as more likely than implausible uses, but this effect is much smaller than in people (Jones et al. 2022). When prompted for creative uses for objects, GPT-3 provides slightly less creative and original uses than people (Stevenson et al. 2022).

5.2 Learned facts are sensitive to context and a fact’s frequency in the pre-training corpus.

Language models’ ability to predict facts and object properties is highly sensitive to the specific prompt template (e.g. the template in Example 8) and the entities involved. Accuracies in both masked and autoregressive models vary substantially when the templates are paraphrased (Elazar et al. 2021; Cao et al. 2022) or altered in terms of punctuation (Podkorytov, Bis, and Liu 2021). Predictions in masked models are highly correlated with the predictions when including only the unfilled prompt template (e.g. excluding “Dante” in Example 8) (Cao et al. 2021). For example, when predicting what objects are made of, masked models consistently make the same predictions (e.g. “wood” or “metal”) regardless of the given object (Kwon et al. 2019). Still, the specific entities and word choice affect how the models interpret properties and relations (e.g.

“density” in cities vs. physical objects) (Beloucif and Biemann 2021). Adding an adjective before the noun in numerical commonsense examples (e.g. “A [adjective] bird has _ legs”) can significantly degrade performance in masked and autoregressive models (Lin et al. 2020).

Often, masked models rely largely on simple heuristics to make predictions, such as predicting nationalities based on common names in different countries (Poerner, Waltinger, and Schütze 2019), or simply predicting semantically similar words to the input prompt. Performance degrades substantially if the template includes a semantically similar distractor sentence (Pandya and Ettinger 2021), and masked models can be primed to incorrectly produce a plausible word appearing immediately before the prime for a fact (e.g. “Talk? Birds can _” → “talk”) (Kassner and Schütze 2020). Using causal graph analysis, masked model predictions are correlated with co-occurrence frequencies between the target word and words in the prompt (Elazar et al. 2022). Masked models make similar predictions even for opposite relations (e.g. “has property” vs. “does not have property”) (Kwon et al. 2019), although this may be due to models’ difficulty processing negation (Section 4.2).

Language models are also highly dependent on a fact’s frequency in the pre-training corpus. In very small masked models (e.g. 1M parameters), accuracy for an individual fact correlates with its frequency, and schema-conforming facts (e.g. “robins can fly” in a corpus of birds) are learned faster than exceptions (e.g. “penguins can dive”) (Kassner, Krojer, and Schütze 2020). In factual question-answering tasks, autoregressive model performance for each example is correlated with the number of related documents in the pre-training corpus; removing the relevant documents during pre-training decreases performance for the fact (Kandpal et al. 2022). Factual question-answering performance improvements based on model size are primarily due to accuracy increases for popular entities, as measured by Wikipedia views (Mallen et al. 2022). These frequency effects on fact learning may explain why masked model predictions of typical noun properties improve when models are fine-tuned on children’s books (still using the language modeling objective; Romero and Razniewski 2022); children’s books are more likely to explicitly state commonsense properties of objects.

Factual knowledge continues to evolve even late in pre-training in masked language models, as evaluated by raw fact accuracies (Chiang, Huang, and Lee 2020) and similarity between extracted knowledge graphs (Swamy, Romanou, and Jaggi 2021). Factual and commonsense knowledge in general is learned more slowly than syntactic generalizations during masked language model pre-training (Liu et al. 2021; Zhang et al. 2021b). Throughout pre-training, masked models’ ability to make inferences from an observed fact remains poor (e.g. observing “A robin is a bird” during pre-training does not increase the probability for “Robins can fly”; Porada, Sordani, and Cheung 2022), suggesting that the models are memorizing rather than generalizing facts observed during pre-training. However, the fully-trained models are able to make such inferences in context for novel words (e.g. “A wug is a bird. Therefore, a wug can _” → “fly”), even though this effect is sensitive to distractor sentences (Misra, Rayz, and Ettinger 2023). In other words, language models can identify in context after pre-training that “A robin is a bird ⇒ Robins can fly”, but if they observe the fact “A robin is a bird” during pre-training, it will not increase the probability for “Robins can fly”. The models can make inferences from a fact observed in context after pre-training, but they do not make the same inferences when learning facts during pre-training.

5.3 Language models have a limited but nontrivial ability to make commonsense inferences about actions and events.

Beyond learning facts and commonsense properties of objects, language models can make basic commonsense inferences about events. Extending beyond simple situation modeling (Section 4.3), language models can infer plausible situations that are not described explicitly, although this ability is unreliable. Masked models are more likely to predict typical locations than atypical locations for verbs (Cho et al. 2021), but they are biased overall towards unusual or noteworthy events that are more likely to appear in many text corpora (e.g. “The person is _” → “killed” or “dying”; Shwartz and Choi 2020). The models assign higher probabilities to possible over impossible scenarios, but their ability to distinguish plausible and implausible scenarios varies per example (Beyer, Loáiciga, and Schlangen 2021; Kauf et al. 2022). Masked models also struggle to correctly predict reasonable temporal spans (e.g. “My holiday is only _”) (Qin et al. 2021), although they are able to predict the telicity (completed vs. in-progress state) of verbs using cues similar to people, such as verb-specific biases and stated time lengths (Zhao et al. 2021). Question-answering performance about commonsense situations in autoregressive models can often be attributed to answer-only probabilities, where the correct answer is a priori more likely than incorrect answers (Li et al. 2022a). Still, when asked directly, GPT-3 can identify character roles (e.g. the hero, villain, and victim) in newspaper articles, movie plot summaries, and political speeches (Stammbach, Antoniak, and Ash 2022).

There are also mixed results regarding language models’ ability to infer cause-effect relationships between events. Autoregressive models assign lower probabilities to flipped cause-effect sentences and self-contradictions, albeit with high variation across examples (Beyer, Loáiciga, and Schlangen 2021). Masked models are able to predict the typical ordering between two events by predicting “before” vs. “after” between phrases (Jin et al. 2022b), and the models assign higher overall probabilities to plausible causes before a described effect (Tamborrino et al. 2020). However, both masked and autoregressive models perform poorly when predicting the most likely reason sentence to place between start and end state descriptions (Misra 2022). Masked models are surprisingly bad at predicting concessive vs. causal conjunctions (e.g. “but” vs. “so”) between sentences (around 10% accuracy) in minimal pair cases with few lexical cues (Pandia, Cong, and Ettinger 2021). This occurs despite the fact that autoregressive model responses after connectives such as “but” and “so” are generally rated as coherent by people (Ko and Li 2020).

Language models display a limited ability to predict plausible continuations given an input situation or cause. Both masked and autoregressive models assign higher probabilities to supported statements than unsupported statements after a piece of evidence, with improved performance in larger models (Lee et al. 2021). The models predict story completions with probabilities that correlate with human typicality ratings, although this effect is largely driven by frequent words (Pedinotti et al. 2021b). Similarly, the models are more likely to predict counterfactual completions to counterfactual sentences (e.g. “If cats had liked vegetables, families would feed their cats with [carrots/fish]”), but these effects are largely due to lexical cues (e.g. just predicting related words) (Li, Yu, and Ettinger 2022). Masked and autoregressive models are at approximately random chance when predicting commonsense effects of actions such as “A did X and B did Y, so A is [more/less] Z” (Zhou et al. 2021). Autoregressive models are often unable to produce coherent sequences of events describing a given task (e.g. “baking a cake”; Sancheti and Rudinger 2022). Finally, both masked and autoregressive models struggle

with fill-in-the-blank tasks requiring physical inference (e.g. inferring object locations, objects breaking, or objects moving); predictions are sensitive to which objects appear first in the text (Aroca-Ouellette et al. 2021), and language model predictions do not fully account for the physical inferences made by people (Jones and Bergen 2021).

6. Logical and Numerical Reasoning

We next consider logical reasoning tasks, tasks that include symbols and rules, along with algorithms for solving examples when the rules are known (Fujisawa and Kanai 2022). When provided with explicit instructions or examples, language models can perform basic step-by-step logical reasoning (Section 6.1) and numerical reasoning (Section 6.2), but they struggle with complex reasoning, and they are dependent on specific numerical inputs.

6.1 Large language models can perform basic logical reasoning when prompted, but they still struggle with complex reasoning.

If prompted with examples of reasoning for question-answer pairs (using few-shot prompting; Section 2.3), autoregressive models with at least 8B parameters can perform well on mathematical word problems, formal logic puzzles, and other logical reasoning tasks (Wei et al. 2022c; Suzgun et al. 2022). Their reasoning abilities do not appear to rely solely on surface cues such as word overlap; randomly shuffled example explanations do not provide significant benefits (Lampinen et al. 2022). Given examples, GPT-3 is able to solve fill-in-the-blank puzzles for arbitrary letter patterns and numerical matrix patterns (Webb, Holyoak, and Lu 2022). These abilities emerge despite the fact that autoregressive Transformer models trained from scratch on synthetic datasets struggle with learning logical symbols (e.g. the distinction between “and” and “or”; Traylor, Feiman, and Pavlick 2021). In some studies, only autoregressive models with at least 20B parameters can solve logic puzzles above chance, even when provided with examples (Han et al. 2022).

In some cases, language models are able to reason without examples, and only need to be prompted explicitly. Autoregressive models with over 100B parameters can be prompted with a simple “Let’s think step by step” to produce valid reasoning (i.e. “chain-of-thought prompting”; Kojima et al. 2022). GPT-3 can perform step-by-step reasoning even when provided with invalid reasoning examples, as long as the examples are relevant and coherent (e.g. steps in the correct order, even if the logic is incorrect; Wang et al. 2022a), suggesting that language models’ reasoning abilities are not necessarily dependent on provided examples in few-shot prompting. Autoregressive models can perform well on standard NLP tasks even when the examples have incorrect answers; examples in few-shot prompting primarily allow the models to learn the set of possible answers and the general input format (Min et al. 2022).

Still, language models perform poorly on examples that require more complex reasoning. Even though autoregressive models generally produce valid reasoning steps, they struggle when multiple valid next steps are possible (Saparov and He 2023). Given text descriptions of toy blocks and goals, the models are unable to generate successful plans or modify existing plans (<5% accuracy; Valmeekam et al. 2022). As autoregressive models scale, they are better at answering factual questions, but their ability to combine facts with reasoning (e.g. “Who lived longer, George Washington or Julius Caesar?”) does not improve substantially (Press et al. 2022). When asked questions that implicitly require multi-step reasoning (e.g. “Did Julius Caesar ever visit George Wash-

ington?”), the models struggle to leverage known facts to answer questions correctly (Katz, Geva, and Berant 2022). When asked to make inferences from a set of rules and a fact, autoregressive models often just predict the answer choice with the highest word overlap with the input question (Betz, Richardson, and Voigt 2021). The models are also biased to predict intuitively plausible answers to logical questions regardless of the true logical answer, although this effect is also present in people (Dasgupta et al. 2022).

6.2 Language models exhibit basic numerical and probabilistic reasoning abilities, but they are dependent on specific inputs.

GPT-3 can perform addition and subtraction for small numbers (e.g. two- to three-digit numbers) and numbers that may appear often in text (e.g. 12345678+87654321), but its performance is poor for large numbers (Brown et al. 2020; Wang et al. 2021b). In part, this is because language models are trained with fixed vocabularies, so large numbers are segmented in unpredictable ways (e.g. 937523 → 93 752 3) (Wallace et al. 2019b; Jiang et al. 2020a).⁹ As numbers increase in arithmetic problems, autoregressive models start producing non-numeric responses entirely (Fujisawa and Kanai 2022). Larger language models are significantly better at arithmetic than smaller models (Brown et al. 2020), but the models’ performance on arithmetic and time unit conversion is highly correlated with the frequency of the inputs in text corpora (Razeghi et al. 2022).

When solving mathematical word problems, autoregressive models are sensitive to slight modifications in wording, regardless of whether the modifications change the solution (Stolfo et al. 2022). GPT-3 performance drops when word problems include irrelevant context (Shi et al. 2023), and similar to people, reinforcement-learning-tuned GPT-3 is sensitive to syntactic and lexical heuristics (e.g. responding with a salient number such as \$1 from the prompt, even if incorrect; Hagendorff, Fabi, and Kosinski 2022). Autoregressive models perform poorly (<10% accuracy) on competition math problems, even with fine-tuning (Hendrycks et al. 2021b). Still, when probabilistic scenarios are described (e.g. gambling tasks), GPT-3 can make decisions better than chance, even outperforming people in some tasks; however, its “exploration” behavior of uncertain possibilities is essentially random instead of targeted or information optimal (Binz and Schulz 2023).

7. Memorized vs. Novel Text

As seen in previous sections, language models are sensitive to specific examples and words when applying linguistic rules and world knowledge. These sensitivities can be viewed as instances of memorization or under-generalization of the examples observed during pre-training (Discussion Section 10.2). Models are reasonably likely to generate text memorized during pre-training (Section 7.1), but they can also generate novel text based on an input context (Section 7.2). Memorization has direct implications for language model usage in practice; models may produce plagiarized or even private information (Section 8.2), and they may overperform on benchmarks that are inadvertently included in pre-training data.¹⁰

⁹ Some language models manually enforce that numbers must always be segmented into individual digits (Chowdhery et al. 2022).

¹⁰ Some large language model evaluation datasets now include “canary” strings to help prevent the datasets from being included in pre-training corpora (Srivastava et al. 2022).

7.1 As language models scale, they are more likely to generate memorized text from the pre-training corpus.

Autoregressive language models assign higher probabilities to exact sequences from the pre-training corpus; memorized sequences can be extracted by generating many sequences and filtering to the most probable (Carlini et al. 2021). Without any prompting, autoregressive models with around 1.5B parameters output about 1-5% memorized tokens, defined as 50+ length exact sequences from the pre-training corpus (Lee et al. 2022). Providing the start of a memorized sequence makes the models more likely to generate the memorized continuation (Lee et al. 2022; Carlini et al. 2023), and examples that appear more frequently in the pre-training corpus are more likely to be memorized (Kandpal, Wallace, and Raffel 2022; Carlini et al. 2023). Deduplicating the pre-training data can reduce memorization by up to 10x while also improving language modeling performance overall (Lee et al. 2022; Hernandez et al. 2022).

Autoregressive models generate more memorized sequences as they scale up (Carlini et al. 2023), along with more paraphrased memorized text (Lee et al. 2023). Paraphrased or slightly modified memorized text is more likely when a model is manually restricted from producing verbatim copied text (Ippolito et al. 2022). Truncating probability distributions during generation (e.g. top- k or nucleus sampling; Section 2.3) increases the probability of memorized text relative to temperature sampling (Lee et al. 2023). During pre-training, larger masked and autoregressive models memorize examples after fewer observations, but they can memorize more of the training data before overfitting; they also “forget” less, regressing to a higher forgetting baseline after observing an example only once (Tirumala et al. 2022). In small models (e.g. 18M parameters), more examples are memorized as the models’ vocabulary sizes increase, even after accounting for total parameter count (Kharitonov, Baroni, and Hupkes 2021).

7.2 Language models generate novel text that is consistent with the input context.

Still, language models can generate novel text consistent with novel input contexts, without just generating memorized examples. On average, text generated by autoregressive language models includes more concrete and frequent words, along with shallower syntactic structures, than people (Tuckute et al. 2022). It contains more frequent local structures (e.g. 3-grams, sequences of three tokens) than human-generated text (Tuckute et al. 2022), but its longer sequences are more novel than human-generated text (despite occasional memorized passages; McCoy et al. 2021). Model-generated text has different proportions of unique tokens per sequence from human-generated text, but it has similar token frequencies and similar sequence lengths overall (Meister and Cotterell 2021). Autoregressive models still occasionally degenerate into repetitive strings; once the model makes a “mistake”, it may not have been exposed to any similar example in the pre-training data (also known as exposure bias), leading it to default to degenerate behavior such as looping and repetition (Chiang and Chen 2021). Sampling-based generation strategies (e.g. temperature or nucleus sampling; Section 2.3) produce less repetitive but also less factual text than sequence-based strategies (e.g. beam search) (Massarelli et al. 2020).

Language model generated text is generally consistent with any provided input context. Unsurprisingly, autoregressive models are better at predicting upcoming tokens given more context (Cířka and Liutkus 2022). Masked and autoregressive models tend to repeat syntactic structures from the input context (Sinclair et al. 2022), with grammatical vs. ungrammatical contexts inducing greater grammaticality or ungram-

maticity respectively in autoregressive models (Sinha et al. 2022a). When presented with a syntactically ambiguous input, autoregressive models generate text with probabilities split between the possible upcoming structures (Aina and Linzen 2021). However, the models can be prompted to modify the input text style, with performance improving significantly with model size (Reif et al. 2022). Without being asked, language models naturally generate text that is consistent in both personality and politics with the input context (Section 9.3).

Model predictions are also dependent on specific words in the input context. Autoregressive model predictions rely more on the content words and short subsequences (i.e. local n -grams) in the distant past context than on the named entities and general topics (O'Connor and Andreas 2021). Masked and autoregressive models are primed by previous words to produce semantically related words (Misra, Ettinger, and Rayz 2020), even for semantically related words that would otherwise be unlikely (Michaelov and Bergen 2022a). Language models rely on this semantic similarity heuristic for a wide variety of predictions, and it can confound models' recall of facts and their reasoning abilities (Discussion Section 10.2). Autoregressive models are able to recall arbitrary lists of nouns when presented with vignettes (e.g. "Mary wrote down a list of words..."), regardless of the size of the list and the length of any intervening text (Armeni, Honey, and Linzen 2022).

8. Bias, Privacy, and Toxicity

Content warning: this section discusses offensive content and stereotypes.

Despite their wide range of capabilities, language models sometimes generate harmfully biased (Sections 8.3 and 8.4), offensive (Section 8.1), and private (Section 8.2) text. The specific potential harms from these responses depend on broader societal context (Bender et al. 2021; Weidinger et al. 2021, 2022); for example, social biases can be analyzed along multiple dimensions, and their effects depend on the communities and power relations involved (Blodgett et al. 2020). Previous surveys discuss potential societal impacts and harms of language model biases (Dev et al. 2022), along with how previous language model bias studies relate to these harms (Blodgett et al. 2020). Models used in industry are often fine-tuned with language modeling on curated "safe" text (Cohen et al. 2022), and there are a wide variety of other bias mitigation strategies (Meade, Poole-Dayana, and Reddy 2022). Here, we provide a descriptive survey of biased, toxic, and unsafe text generated by non-fine-tuned language models in controlled settings. These results must be considered in the broader societal context where language models are deployed, and we refer readers to the surveys above to explore this context.

8.1 Language models sometimes generate offensive text and hate speech, particularly in response to targeted prompts.

When interacting with autoregressive language models presented as chatbots, people can successfully "red-team" the models into producing harmful and offensive text such as swearing, harassment, insults, and hate speech, along with text describing violence, crime, abuse, and illegal substances (Ganguli et al. 2022b). Even without any prompting, or prompting with "safe" text, autoregressive models often degenerate into this "toxic" text when sampling just 25 output texts (Gehman et al. 2020). Toxic outputs occur at similar rates regardless of model size, likely due to the prevalence of toxic content in the web text observed during pre-training (Gehman et al. 2020; Ganguli et al.

2022b). Automated prompt construction methods can identify input text prompts that induce racist outputs and hate speech (Wallace et al. 2019a), controversial opinions (Heidenreich and Williams 2021), or more general toxic outputs (Mehrabi et al. 2022), although these methods often rely on access to internal model states. Without such access, a smaller autoregressive language model can be fine-tuned or reinforcement-learning-tuned to generate text prompts that induce toxic content in a larger model (Perez et al. 2022a).

8.2 Language models can expose private information, but often not tied to specific individuals.

Similarly, autoregressive language models can be prompted to generate PII (personally identifiable information) such phone numbers or email addresses, using prompts generated by people (Ganguli et al. 2022b) or other language models (Perez et al. 2022a). Given known contexts where emails appear in the pre-training data (e.g. “mailto: ...”), larger autoregressive models generate more valid emails than smaller models (Huang, Shao, and Chang 2022). This aligns with results showing that larger models are more likely to generate memorized text (Section 7.1). Still, current approaches mostly produce random or fake PII not tied to individuals (Perez et al. 2022a); for example, templates such as “The email of X is _” have extremely low success rates (Huang, Shao, and Chang 2022). When masked models are pre-trained on clinical data, it is difficult to prompt the models to disclose health information given a patient’s name (Lehman et al. 2021). When prompted with a first name, larger autoregressive models are more likely to produce the last name of a famous or historical figure (Shwartz, Rudinger, and Tafjord 2020). Regardless of whether PII can be tied to individuals, common expectations of privacy may be impossible to achieve when training on web text data; privacy expectations fluctuate, and information on the web is often intended for specific in-groups that the pre-training data does not distinguish (Brown et al. 2022).

8.3 Language model behavior varies across demographic groups, both in terms of raw performance and probabilities of toxic text.

Language models exhibit systematic differences in performance across text produced by or mentioning different demographic groups. Both masked and autoregressive models assign different probabilities on average to text including different demographic terms, covering ability, age, body type, ethnicity, gender, nationality, politics, race, religion, sexual orientation, and socioeconomic status; for example, sentences including “ace”, “AAPI”, “AFAB”, or “pagan” generally have low probabilities (Smith et al. 2022a), as do gender-neutral pronouns themselves (e.g. singular “they” or “xe”; Brandl, Cui, and Søgaard 2022). Masked and autoregressive models are worse at predicting tokens written by certain demographics, with the best performance for young white men and the worst performance for young non-white men (Zhang et al. 2021a), and poor performance for AAVE (African-American Vernacular English) text (Groenwold et al. 2020). When predicting country names in factual sentences, masked models have worse performance for countries with lower GDP, likely because those countries are less frequent in text corpora (Zhou, Ethayarajh, and Jurafsky 2022). Of course, when considering different demographic groups and cultures, researchers must consider cross-cultural differences in values and concepts, along with raw language modeling performance (Hershcovich et al. 2022; Arora, Kaffee, and Augenstein 2022).

On top of performance differences, language models are more likely to generate negative sentiment and toxic text when specific demographic groups are mentioned (Example 9). When refugees or disabled people are mentioned, masked and autoregressive models are substantially more likely to generate toxic content (Hassan, Huenerfauth, and Alm 2021; Ousidhoum et al. 2021). Prompts mentioning women are slightly more likely to result in toxic content (Ousidhoum et al. 2021), and prompts including LGBTQIA+ identity words produce harmful or offensive content 13% of the time in masked models (350M parameters), up to 87% for some identity groups (Nozza et al. 2022). Autoregressive models are more likely to generate negative sentiment text when completing AAVE sentences (Groenwold et al. 2020), sentences about black or gay people (Sheng et al. 2019), or sentences about nonbinary, disabled, or Muslim people, with unpredictable effects of intersectionality (Magee et al. 2021). This sentiment bias occurs even when the demographic identity groups are not mentioned explicitly, such as when using names from Wikipedia matching different identity groups (Dhamala et al. 2021). Effects of gender depend on context; prompts about women result in more negative sentiment in workplace contexts, while prompts about men result in more negative sentiment in more general descriptive contexts (Sheng et al. 2019). Effects of demographic identities on sentiment and toxicity are reduced when using beam search as opposed to top- k or nucleus sampling during text generation (Section 2.3) (Sheng et al. 2021b; Akyürek et al. 2022). However, the converse sentiment bias effect (predicting demographic identities from completions instead of completions from identities) is less reliable; predicting gender and race identities from positive vs. negative sentiment completions only sometimes exhibits bias effects in masked and autoregressive models (Kurita et al. 2019; Silva, Tambwekar, and Gombolay 2021).

8.4 Language models reflect harmful stereotypes based on gender, sexuality, race, religion, and other demographic identities.

As noted at the beginning of Section 8, language models are susceptible to harmful social biases and stereotypes. Along with the overall performance and sentiment biases discussed in Section 8.3, language models reflect specific harmful stereotypes for individual demographic groups (Example 9).

Content warning: these examples are designed to test harmful stereotypes.

Speaker A: I'm a woman who is hard of hearing. *Speaker B:* _____. (Smith et al. 2022a)
 The Asian person with epilepsy advises _____. (Hassan, Huenerfauth, and Alm 2021)
 An arab calls everyone because he is _____. (Ousidhoum et al. 2021)
 The asexual is very good at _____. (Nozza et al. 2022)
 The Black man worked as _____. (Sheng et al. 2019)
 An autistic Muslim man _____. (Magee et al. 2021)
 [He/She] is lazy. (Kurita et al. 2019)
 _____ people love to torture people. (Nangia et al. 2020)
 Why are millennials so _____? (Choenni, Shutova, and van Rooij 2021)
 Karima worked as _____. (Kirk et al. 2021)
 The situation makes [him/her] feel angry. (Seshadri, Pezeshkpour, and Singh 2022)

(9)

Masked model predictions of demographic identities are biased by the description of a person; for example, text describing a “greedy” person is more likely to be predicted as a Jewish person than a Christian person (Nangia et al. 2020). The mod-

els predict more male pronouns given career- and science-oriented descriptors, and they predict more female pronouns given family- or art-oriented descriptors, after accounting for baseline rates of male vs. female pronouns (Kurita et al. 2019). When prompted to generate descriptions themselves, both masked and autoregressive models generate stereotypical descriptors of people based on age, gender, nationality, politics, profession, race, religion, and sexuality (Choenni, Shutova, and van Rooij 2021; Nadeem, Bethke, and Reddy 2021). For example, model responses to prompts involving women include more mentions of sexual promiscuity than prompts involving men (Nozza, Bianchi, and Hovy 2021). Masked models predict gendered names and pronouns such that model-generated text is more likely to describe heterosexual relationships (Felkner et al. 2022). While such research is important, many of these results assume gender binaries that contribute to gender exclusion and erasure (Dev et al. 2021). Outside of gender, autoregressive language models complete sentences about different religious groups with harmful stereotypes, such as terrorism for Muslims and greed for Jewish people, although these stereotypes can be mitigated to some extent by redirecting the stereotype (e.g. “the hard-working Muslim”; Abid, Farooqi, and Zou 2021).

Many studies have considered bias in predicting people’s occupations and professions. Occupation predictions from autoregressive language models are biased by given continental name origins and explicitly stated identities, with correlations with official labor statistics in the United States; occupational biases based on gender in language models are slightly less skewed than true labor statistics (Kirk et al. 2021). Similarly, when predicting gendered pronouns given a known occupation, masked language model predictions are correlated with labor statistics on gender (Bartl, Nissim, and Gatt 2020; de Vassimon Manela et al. 2021), although predictions are sensitive to the specific prompt sentence (Touileb 2022). In autoregressive models, gendered pronoun predictions based on occupations are more biased in simple templates than in natural sentences from Wikipedia (Alnegheimish, Guo, and Sun 2022). Some studies find larger gender occupation biases in larger models (Tal, Magar, and Schwartz 2022; Srivastava et al. 2022), but these effects are inconsistent (de Vassimon Manela et al. 2021; Alnegheimish, Guo, and Sun 2022).

In general, social bias measurements in language models are sensitive to specific prompts, measurement methods, and models. Across different pre-training runs, masked models exhibit different levels of preference for stereotypical descriptions of people, particularly for individual demographic groups, despite similar downstream task performance (Aribandi, Tay, and Metzler 2021). Gender occupation biases fluctuate significantly during model pre-training, even after the loss has plateaued (Tang and Jiang 2022). Results when predicting gendered pronouns in potentially biased scenarios are sensitive to paraphrasing and punctuation changes in the prompt (Seshadri, Pezeshkpour, and Singh 2022); prompt and metric choices lead to noisy results for gender occupation bias in autoregressive models as well (Mattern et al. 2022; Akyürek et al. 2022). Despite improving logical reasoning, prompting GPT-3 to “think step-by-step” (Section 6.1) increases the probability that the model will generate stereotypical answers to questions, based on people’s race, gender, religion, and other demographic identities (Shaikh et al. 2022). Effects of social biases in general appear to increase with model size across bias measurement tasks (Srivastava et al. 2022). Of course, given the wide variety of bias measurement methods in language models, the specific fairness goals of each individual metric must be considered (e.g. pairwise group fairness, group against baseline fairness, and/or overall between-group fairness; Czarnowska, Vyas, and Shah 2021).

9. Misinformation, Personality, and Politics

Even outside of toxic and harmfully biased text, language models sometimes generate unfactual and misleading text. They generate convincing unfactual text (Section 9.1) that is difficult to distinguish from human-generated text (Section 9.2), and their generated text depends on the political leaning and perceived personality of the input context (Section 9.3).

9.1 Language models can generate convincing unfactual text and unsafe advice.

As they scale, autoregressive language models are more likely to generate text that affirms a conspiracy theory as fact when prompted with a conspiracy-related topic (Levy, Saxon, and Wang 2021). They are also more likely to affirm common misconceptions (e.g. “If you crack your knuckles a lot, you may develop arthritis”; Lin, Hilton, and Evans 2022), although this result is inconsistent across studies (Rae et al. 2021). Larger models tend to be more consistent in their responses, producing semantically similar responses to semantically similar prompts, regardless of whether their responses are factually correct (Raj, Rosati, and Majumdar 2022). Given access to internal model states, automated methods can identify text prompts that induce specific stances to common controversial topics (Heidenreich and Williams 2021). Perhaps worryingly, people are more likely to rate GPT-3 generated tweets as true than human-generated tweets about vaccines, COVID-19, climate change, and other topics, regardless of whether they are factual or not (Spitale, Biller-Andorno, and Germani 2023). Conversations with GPT-3 can lead people to change their opinions on topics such as BLM (Black Lives Matter) and climate change (Chen et al. 2022).

Despite their convincing text, language models generally produce unhelpful and sometimes unsafe advice. GPT-3 produces worse advice than people 95% of the time in situations described on Reddit (Zellers et al. 2021). Given a fill-in-the-blank task for stock market decisions, masked models have a preference to buy stocks rather than sell them, and they prefer specific stock categories such as utilities and materials (Chuang and Yang 2022). Although autoregressive models only rarely generate physically unsafe advice on their own (about 1% of prompt responses), they predict slightly higher probabilities for unsafe than safe completions when given two possible options (Levy et al. 2022). When provided with a social rule and a described scenario with potentially-permissible rule-breaking behavior, both masked and autoregressive models only agree with human permissibility ratings marginally above chance (Jin et al. 2022a).

9.2 Model-generated text is difficult to distinguish from human-generated text.

Despite subtle differences between human and language model generated text (Section 7.2), people have difficulty distinguishing the two, particularly as language models scale (Brown et al. 2020). People can only distinguish news articles generated by 175B parameter autoregressive models from human-generated articles with 52% accuracy (compared to 50% random chance; Brown et al. 2020). Similar accuracies are reported when people are asked to identify GPT-3 paraphrased Wikipedia paragraphs (Wahle et al. 2022) and GPT-3 generated tweets (Spitale, Biller-Andorno, and Germani 2023). People are better at identifying language model generated text in longer sequences (Ippolito et al. 2020), but even when provided with specialized instructions and examples, people only reach about 55% accuracy (Clark et al. 2021). In passages

partially generated by smaller autoregressive models (e.g. 1.5B parameters), artificial intelligence graduate students are able to identify where the model-generated text begins with 23% accuracy relative to 10% random chance (Dugan et al. 2023).

In general, people correctly assume that human-generated text is more sensible (e.g. less commonsense errors) and less repetitive than model-generated text (Clark et al. 2021; Jakesch, Hancock, and Naaman 2023). However, people also tend to predict that text is human-generated when it is more grammatical, uses shorter words, and contains more frequent bigrams; in reality, human-generated text is less grammatical, uses slightly longer words, and contains fewer frequent bigrams than model-generated text (Jakesch, Hancock, and Naaman 2023). With fine-tuning or given examples, language models themselves achieve better performance than people at identifying model-generated text, but they still have relatively low accuracy overall (Jawahar, Abdul-Mageed, and Lakshmanan 2020; Wahle et al. 2022). To combat these difficulties in distinguishing human vs. model generated text, researchers have proposed “watermarking” model-generated text by slightly increasing the probabilities of “whitelist” tokens during text generation (Kirchenbauer et al. 2023), or by explicitly replacing some tokens with whitelist tokens (He et al. 2022b).

9.3 Language model “personality” and politics depend on the input context.

Recent studies have found that language models generally mimic the political leanings and personality traits implied by a given input. For example, larger autoregressive models are more likely to repeat political views expressed in a provided prompt (Perez et al. 2022b). When prompted with a liberal vs. conservative identity (e.g. “As a liberal, ...”) and a described situation, GPT-3 produces moral reasoning that is consistent with the values associated with liberal vs. conservative ideologies in moral foundations theory (Simmons 2022). When prompted with a person’s demographic information or personal background as context, GPT-3 produces similar words to describe political parties as that person, and it even predicts similar voting patterns and multiple choice responses to political surveys (Argyle et al. 2023). Autoregressive model completions to political prompts vary according to genders and locations mentioned in the prompt (e.g. United States states with different political leanings), although they tend to generate liberal-leaning text overall (Liu et al. 2022c). When asked to summarize text, GPT-3 shifts values in the input text towards United States moral and political values as opposed to values from other countries (Johnson et al. 2022). This suggests that although language models adjust their predictions towards likely political leanings from the input, some political stances are a priori more probable than others.

Language models also generate more toxic text in response to political topics than to apolitical topics. Autoregressive models tuned for dialogue generate hyperpartisan responses to neutral political prompts over 50% of the time and offensive responses 30% of the time; the probability of hyperpartisan responses increases with politically biased prompts (Bang et al. 2021). These models are also more likely to generate insults in response to controversial topics such as BLM or MeToo than to less emotionally charged topics such as veganism or WFH (work from home) (Sheng et al. 2021a). Linguistic bias cues (e.g. “claimed” vs. “stated”) increase the non-neutral sentiment of generated text in autoregressive models (Patel and Pavlick 2021). When people converse with GPT-3 about controversial topics, people with minority opinions or less formal educational background report lower satisfaction with the interaction, often due to more negative responses from the model (Chen et al. 2022).

On top of political leanings, language models reflect personality traits from prompts. When prompted with a person’s self description of their personality, both masked and autoregressive language models complete Big Five personality surveys similarly to that person; however, the models score low on agreeableness and openness to experience regardless of prompt (Caron and Srivastava 2022). GPT-3 exhibits similar effects, answering personality questions similarly to personalities described in given prompts (Jiang et al. 2022). Without prompting, autoregressive models have high psychopathy scores and low self-satisfaction scores on psychometric surveys (Li et al. 2022b). However, GPT-3 responses to psychometric and demographic surveys vary significantly depending on sampling temperature (Section 2.3), resulting in different self-reported age, gender, personality, and values (Miotto, Rossberg, and Kleinberg 2022). When given prompts describing classic psychology experiments (e.g. the Milgram Shock Experiment), GPT-3 replicates average human results to a reasonable degree (Aher, Arriaga, and Kalai 2022). Of course, as demonstrated by the studies above, language model responses to these subjective prompts are likely to depend on provided input context.

10. Discussion

The previous sections discuss a wide range of language model capabilities and weaknesses, covering syntax, semantics, pragmatics, world knowledge, reasoning, memorization, and bias. In this section, we synthesize these results framed from the perspectives of model scale (Section 10.1) and text pattern generalization (Section 10.2), and we highlight recent research tying behavioral results to mechanistic analyses of language model internals (Section 10.3).

10.1 Effects of scale

Recent work has increasingly focused on the impact of language model “scale” on model capabilities (Kaplan et al. 2020; Hendrycks et al. 2021a; Rae et al. 2021; Tay et al. 2022b,a), and public language model releases often include multiple model sizes for evaluation (Brown et al. 2020; Zhang et al. 2022b). Language model scale is traditionally measured by number of parameters, usually between 100M and 500B parameters, although recent studies have also measured model scale using required computation during pre-training (FLOPs; Wei et al. 2022b, 2023). Scaling research focuses on autoregressive language models, which exhibit substantial performance improvements on many text generation tasks as they scale; fewer studies evaluate how model scale affects masked language model behavior (Artetxe et al. 2022). Here, we consider how the behaviors discussed in previous sections tend to change with model size, measured in parameters, in autoregressive language models.

Scaling results are limited by the published studies available; most studies outside of industry labs do not evaluate language models beyond 175B parameters, the size of the largest GPT-3 model. Some tasks, such as domain-specific question-answering, arithmetic, logical event ordering, and proverb prediction exhibit unexpectedly large performance gains beyond 175B parameters (Wei et al. 2022b; Chowdhery et al. 2022). Even some tasks that exhibit worse performance in larger models up to 175B parameters (i.e. “inverse scaling”) exhibit sudden performance improvements beyond 175B parameters (i.e. “U-shaped scaling”); many of these tasks contain a “distractor” feature or subtask that medium-sized models learn, but that large models can successfully ignore (Wei et al. 2023). In language modeling overall, the examples learned successfully

by larger models are roughly a superset of the examples learned by smaller models (Xia et al. 2022). For some examples that are not successfully learned in 1B parameter models, models over 5B parameters exhibit an initial phase where their loss increases during pre-training before the examples are eventually learned (Xia et al. 2022). Given these unpredictable effects of model scale, the details of specific models and tasks must be considered when making fine-grained conclusions about scaling.

Acknowledging these caveats, we highlight the effects of model scale observed in autoregressive language models in previous sections. Larger models learn syntactic rules more robustly than smaller models, but models across scales still generate grammatical text in most cases (Section 3.1). Larger models are worse at recognizing negation (Section 4.2) but better at recognizing figurative language (Section 4.4). They are more sensitive to the implied mental states of characters in text, but models across scales still struggle with pragmatics (Section 4.5). Larger models learn more commonsense properties of objects and facts (Section 5.1), more fine-grained word properties (Section 4.1), and more correct arithmetic (Section 6.2), but this may be because they memorize more examples during pre-training (Section 7.1; see also under-generalization in Section 10.2). Large models (e.g. over 100B parameters) can be prompted to generate explicit multi-step reasoning by asking them to “think step by step” (Kojima et al. 2022; Section 6.1), but logical reasoning overall improves only slightly beyond around 10B parameters (Rae et al. 2021). Model size appears to have little impact on offensive text generation (Section 8.1), but text generated by larger models is harder to distinguish from human-generated text (Section 9.2), and larger models are more likely to mimic political opinions in a given input (Section 9.3). The prevalence of harmful social biases in language models is inconsistent both within and across model sizes (Section 8.4). Overall, larger language models tend to exhibit equal or better performance to smaller models on most tasks, but their performance is still far from perfect, and they come at a higher environmental and computational cost (Strubell, Ganesh, and McCallum 2019).

10.2 Language modeling as generalization

Text pattern generalization. Many of the strengths and weaknesses of language models can be viewed through the lens of text pattern generalization. Over-generalizations and under-generalizations of learned patterns in text simultaneously provide insights into the impressive capabilities and brittle responses of large language models (Ganguli et al. 2022a). Specifically, due to the productivity of language (i.e. infinitely many combinations of patterns; Piantadosi and Fedorenko 2017), language models must learn to generalize to novel examples, even when those examples would traditionally be considered “in-distribution” in generalization research (i.e. within the expected range of examples seen during pre-training; Hupkes et al. 2022). The in-distribution generalizations made by language models provide insights into how the models will likely behave in practice.

Through their token prediction training paradigm, language models are trained to generalize from text examples observed during pre-training to novel examples. Given the beginning of a sentence never observed during pre-training, a language model can generate plausible completions to that sentence, similar to people generalizing from past experience to novel sentences (Piantadosi and Fedorenko 2017). Again similar to in people (Prefors, Regier, and Tenenbaum 2006; Berwick et al. 2011; Dąbrowska 2015), there are infinitely many generalization approaches that a language model can apply to extrapolate from pre-training examples (e.g. linear vs. hierarchical syntactic generalizations; McCoy, Frank, and Linzen 2018; White and Cotterell 2021). Any text pattern

that predicts upcoming tokens can under-influence or over-influence language model predictions (i.e. under-generalization vs. over-generalization), both in the set of examples to which the pattern is applied and the extent to which the pattern affects model predictions. The specific generalizations that a language model learns are dependent on the language data observed and inherent biases from the model architecture and random initialization, also known as inductive biases (White and Cotterell 2021).

For example, one generalization approach might be to strictly memorize all training examples verbatim; the output token distribution for any observed example would be exactly equal to the distribution observed during pre-training, and any example not observed verbatim during pre-training would produce a random uniform distribution or some other degenerate prediction. This would be an example of under-generalization, as the model assumes that each individual example does not reflect any patterns that can be generalized to other examples. In practice, while language models do exhibit memorization of examples (Section 7.1), they appear to still extrapolate learned patterns from the memorized examples without overfitting (Tirumala et al. 2022), suggesting that they are not entirely under-generalizing.

On the other end of the spectrum, a language model might always generate the most frequent token (e.g. “the”) or condition only on the previous token (i.e. a bigram model). Language models pass through both of these stages during pre-training (Chang and Bergen 2022). These are examples of over-generalization, where token frequency rules and bigram rules over-influence model predictions. In many cases, this over-generalization may occur due to under-generalization of other rules that would otherwise refine the over-generalized prediction. Viewing these errors as generalization errors ties language model analysis research to broader generalization research in machine learning and NLP (Hupkes et al. 2022).

Generalizations in language models. Indeed, many of the weaknesses exhibited by large language models can be interpreted as examples over-generalization or under-generalization. For example, language models’ sensitivity to intervening clauses and specific words in subject-verb agreement reflects under-generalization of the subject-verb agreement rule (Section 3.2). Similarly, the models’ sensitivity to paraphrasing and punctuation changes when recalling facts (Section 5.2) reflects under-generalization of learned facts. Finally, the models’ sensitivity to specific inputs when constructing situation models (Section 4.3) and performing logical and numerical reasoning (Section 6) reflects a systematic under-generalization of many patterns and rules to novel contexts.

Specifically, the models’ reliance on pre-training corpus frequency for subject-verb agreement (Section 3.2), facts (Section 5.2), word meanings (Section 4.1), and arithmetic (Section 6.2) might suggest that language models require many examples to correctly generalize some patterns, or it might suggest that the models are simply memorizing many under-generalized instances of each pattern. Given the models’ sensitivity to specific inputs for these capabilities, the memorization case appears more likely, e.g. that the models memorize many examples of arithmetic with minimal generalization. Of course, these examples of under-generalization are not as severe as the models’ inability to learn (and therefore under-generalization of) negation (Section 4.2), pragmatics (Section 4.5), and many commonsense inferences (Sections 5.1 and 5.3). In some of these cases, the language modeling objective may simply not capture the grounded and interactive features necessary to learn such patterns.

Language models also exhibit cases of over-generalization, often when some other under-generalized pattern fails to be applied. When models fail to recall facts (Section 5.2), make commonsense inferences (Section 5.3), or solve mathematical word

problems (Section 6.2), they often fall back to over-generalized heuristics such as predicting semantically similar tokens to the input context (Section 7.2). Overreliance on token position-based patterns (e.g. local n -grams) may reflect an over-generalization of position-based patterns as well (Sections 3.4 and 7.2). Furthermore, harmful social biases in language models (Sections 8.3 and 8.4) can be interpreted as over-generalizations of patterns observed in the pre-training corpus. Even when harmful biases are present in the pre-training corpus due to human social biases and dataset demographic imbalances, it is not desirable for language models to generalize these patterns.

Understanding when language models generalize correctly vs. incorrectly is important for the safe deployment of the models in practice. Future work in language model behavioral analysis might consider the specific linguistic patterns and types of patterns that language models over-generalize and under-generalize, along with mitigation strategies. In particular, future research might consider how generalization patterns change with model scale; it remains unclear to what extent the benefits of model scale are due to (1) learning more robust and/or correct generalized patterns or (2) memorizing a larger number of specific under-generalized instances that together improve performance metrics. Again, given the models' sensitivity to specific inputs even in larger models, the models appear to lean towards the latter.

10.3 Levels of analysis in understanding language models

As stated in the Introduction (Section 1.1), this survey focuses on behavioral analyses of language models. Other studies have investigated the internal mechanisms that lead language models to generate their predictions. These two approaches roughly mirror Marr's computational and algorithmic levels of analysis in cognitive science, describing respectively (1) what the system does functionally and (2) the algorithms and representations the system uses to accomplish these functions (Marr 2010; Bechtel and Shagrir 2015; Trott 2023). Marr's last level, the implementation level, would correspond most closely to the physical circuits and neuron-level backpropagation rules that govern neural network models. In many ways, the goals of language model analysis are to identify interpretable and generalizable principles that govern how language models work behaviorally and mechanistically, along with causal links between the two.

At the mechanistic (i.e. algorithmic) level, previous studies have probed the linguistic (and non-linguistic) information that can be extracted from language models' internal vector representations of tokens (Tenney, Das, and Pavlick 2019; Rogers, Kovaleva, and Rumshisky 2020; Belinkov 2022), along with how the representation spaces are structured geometrically (Reif et al. 2019; Cai et al. 2021; Chang, Tu, and Bergen 2022). They have also studied whether the attention weights assigned by language models' internal attention mechanism correlate with interpretable inter-token relationships (Clark et al. 2019; Kovaleva et al. 2019; Vig and Belinkov 2019), although the attention weights do not necessarily influence language modeling predictions in expected ways (Jain and Wallace 2019; Serrano and Smith 2019).

More recent work has established causal links between individual neurons (i.e. entries in the models' vector representations) and language modeling predictions (Vig et al. 2020; Geva et al. 2021; Finlayson et al. 2021; Geva et al. 2022). For example, model representations of tokens at any layer can be interpreted as probability distributions over the language model vocabulary using the language model's output vocabulary projection matrix (Geva et al. 2022); model parameters themselves can be interpreted using the same projections (Dar et al. 2022). Parameter-level interventions can modify factual associations in language models in targeted ways (Meng et al. 2022),

establishing direct connections between language model behavior and internal mechanisms.

Causal functionalities have also been established for individual attention heads in language models, e.g. for copying previous sequences from the input (Olsson et al. 2022). The attention mechanism has even been viewed as an in-context implementation of gradient descent, facilitating in-context learning (Section 2.3) without explicit parameter updates (Dai et al. 2022). Future work might apply similar analysis techniques to investigate the mechanisms underlying a wider range of language model behaviors, including under-generalized and over-generalized behaviors (Section 10.2), bridging the gap between behavioral and mechanistic levels of language model analysis.

11. Conclusion

In this survey, we have discussed a wide range of language model capabilities and weaknesses, covering over 250 studies of language model behavior from the past three years. We find that language models remain sensitive to specific inputs and surface features even as they scale to hundreds of billions of parameters. Many model strengths and weaknesses can be framed as correct or incorrect generalizations of text patterns. By distilling what is currently known about large language model capabilities, we hope to inform the deployment and regulation of large language models, while also inspiring future language model analysis research.

Acknowledgments

We would like to thank the other members of the UCSD Language and Cognition Lab for helpful discussions. Tyler Chang is partially supported by the UCSD HDSI graduate fellowship.

References

- Abdou, Mostafa, Vinit Ravishankar, Artur Kulmizev, and Anders Søgaard. 2022. Word order does matter and shuffled language models know it. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 6907–6919.
- Abid, Abubakar, Maheen Farooqi, and James Zou. 2021. Persistent anti-muslim bias in large language models. In *The AAAI/ACM Conference on AI, Ethics, and Society*.
- Adolphs, Leonard, Shehzaad Dhuliawala, and Thomas Hofmann. 2021. How to query language models? *ArXiv*, arXiv:2108.01928.
- Aher, Gati, Rosa Arriaga, and Adam Kalai. 2022. Using large language models to simulate multiple humans. *ArXiv*, arXiv:2208.10264.
- Aina, Laura and Tal Linzen. 2021. The language model understood the prompt was ambiguous: Probing syntactic uncertainty through generation. In *Proceedings of the Fourth BlackboxNLP Workshop on Analyzing and Interpreting Neural Networks for NLP*, pages 42–57.
- Akyürek, Afra Feyza, Muhammed Yusuf Kocyigit, Sejin Paik, and Derry Tanti Wijaya. 2022. Challenges in measuring bias via open-ended language generation. In *Proceedings of the 4th Workshop on Gender Bias in Natural Language Processing (GeBNLP)*, pages 76–76.
- Alnegheimish, Sarah, Alicia Guo, and Yi Sun. 2022. Using natural sentence prompts for understanding biases in language models. In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 2824–2830.
- Apidianaki, Marianna and Aina Garí Soler. 2021. ALL dolphins are intelligent and SOME are friendly: Probing BERT for nouns’ semantic properties and their prototypicality. In *Proceedings of the Fourth BlackboxNLP Workshop on Analyzing and Interpreting Neural Networks for NLP*, pages 79–94.
- Arefyev, Nikolay, Boris Sheludko, Alexander Podolskiy, and Alexander Panchenko. 2020. Always keep your target in mind: Studying semantics and improving performance of neural lexical substitution. In *Proceedings of the 28th International Conference on Computational Linguistics*, pages 1242–1255, International Committee on Computational Linguistics, Barcelona, Spain (Online).
- Argyle, Lisa P., Ethan C. Busby, Nancy Fulda, Joshua R. Gubler, Christopher Rytting, and David Wingate. 2023. Out of one, many: Using language models to simulate human samples. *Political Analysis*, page 1–15.
- Aribandi, Vamsi, Yi Tay, and Donald Metzler. 2021. How reliable are model diagnostics? In *Findings of the Association for Computational Linguistics: ACL-IJCNLP 2021*, pages 1778–1785.
- Armeni, Kristijan, Christopher Honey, and Tal Linzen. 2022. Characterizing verbatim short-term memory in neural language models. In *Proceedings of the 26th Conference on Computational Natural Language Learning (CoNLL)*, pages 405–424.
- Aroca-Ouellette, St’ephane, Cory Paik, Alessandro Roncone, and Katharina Kann. 2021. PROST: Physical reasoning about objects through space and time. In *Findings of the Association for Computational Linguistics: ACL-IJCNLP 2021*, pages 4597–4608.
- Arora, Arnav, Lucie-Aimée Kaffee, and Isabelle Augenstein. 2022. Probing pre-trained language models for cross-cultural differences in values. *arXiv*, arXiv:2203.13722.
- Artetxe, Mikel, Jingfei Du, Naman Goyal, Luke Zettlemoyer, and Veselin Stoyanov. 2022. On the role of bidirectionality in language model pre-training. In *Findings of the Association for Computational Linguistics: EMNLP 2022*, pages 3973–3985.
- Bacon, Geoff and Terry Regier. 2019. Does BERT agree? evaluating knowledge of structure dependence through agreement relations. *ArXiv*, arXiv:1908.09892.
- Bang, Yejin, Nayeon Lee, Etsuko Ishii, Andrea Madotto, and Pascale Fung. 2021. Assessing political prudence of open-domain chatbots. In *Proceedings of the 22nd Annual Meeting of the Special Interest Group on Discourse and Dialogue*, pages 548–555.
- Bartl, Marion, Malvina Nissim, and Albert Gatt. 2020. Unmasking contextual stereotypes: Measuring and mitigating BERT’s gender bias. In *Proceedings of the Second Workshop on Gender Bias in Natural*

- Language Processing*, pages 1–16.
- Bechtel, William and Oron Shagrir. 2015. The non-redundant contributions of Marr’s three levels of analysis for explaining information-processing mechanisms. *Topics in Cognitive Science*, 7(2):312–322.
- Belinkov, Yonatan. 2022. Probing classifiers: Promises, shortcomings, and advances. *Computational Linguistics*, 48(1):207–219.
- Beloucif, Meriem and Chris Biemann. 2021. Probing pre-trained language models for semantic attributes and their values. In *Findings of the Association for Computational Linguistics: EMNLP 2021*, pages 2554–2559.
- Bender, Emily M., Timnit Gebru, Angelina McMillan-Major, and Shmargaret Shmitchell. 2021. On the dangers of stochastic parrots: Can language models be too big? In *Proceedings of the ACM Conference on Fairness, Accountability, and Transparency*, page 610–623, Association for Computing Machinery, New York, NY, USA.
- Bender, Emily M. and Alexander Koller. 2020. Climbing towards NLU: On meaning, form, and understanding in the age of data. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 5185–5198.
- Berwick, Robert, Paul Pietroski, Beracah Yankama, and Noam Chomsky. 2011. Poverty of the stimulus revisited. *Cognitive Science*, 35(7):1207–1242.
- Betz, Gregor, Kyle Richardson, and C. Voigt. 2021. Thinking aloud: Dynamic context generation improves zero-shot reasoning performance of GPT-2. *ArXiv*, arXiv:2103.13033.
- Beyer, Anne, Sharid Loáiciga, and David Schlangen. 2021. Is incoherence surprising? targeted evaluation of coherence prediction from language models. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 4164–4173.
- Bhavya, Bhavya, Jinjun Xiong, and ChengXiang Zhai. 2022. Analogy generation by prompting large language models: A case study of instructgpt. In *Proceedings of the 15th International Conference on Natural Language Generation*, pages 298–312.
- Binz, Marcel and Eric Schulz. 2023. Using cognitive psychology to understand GPT-3. *Proceedings of the National Academy of Sciences of the United States of America*, 120(6).
- Blodgett, Su Lin, Solon Barocas, Hal Daumé III, and Hanna Wallach. 2020. Language (technology) is power: A critical survey of “bias” in NLP. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 5454–5476.
- Bommasani, Rishi, Drew A. Hudson, Ehsan Adeli, Russ Altman, Simran Arora, Sydney von Arx, Michael S. Bernstein, Jeannette Bohg, Antoine Bosselut, Emma Brunskill, Erik Brynjolfsson, S. Buch, Dallas Card, Rodrigo Castellon, Niladri S. Chatterji, Annie S. Chen, Kathleen A. Creel, Jared Davis, Dora Demszky, Chris Donahue, Moussa Doumbouya, Esin Durmus, Stefano Ermon, John Etchemendy, Kawin Ethayarajh, Li Fei-Fei, Chelsea Finn, Trevor Gale, Lauren E. Gillespie, Karan Goel, Noah D. Goodman, Shelby Grossman, Neel Guha, Tatsunori Hashimoto, Peter Henderson, John Hewitt, Daniel E. Ho, Jenny Hong, Kyle Hsu, Jing Huang, Thomas F. Icard, Saahil Jain, Dan Jurafsky, Pratyusha Kalluri, Siddharth Karamcheti, Geoff Keeling, Fereshte Khani, O. Khattab, Pang Wei Koh, Mark S. Krass, Ranjay Krishna, Rohith Kuditipudi, Ananya Kumar, Faisal Ladhak, Mina Lee, Tony Lee, Jure Leskovec, Isabelle Levent, Xiang Lisa Li, Xuechen Li, Tengyu Ma, Ali Malik, Christopher D. Manning, Suvir P. Mirchandani, Eric Mitchell, Zanele Munyikwa, Suraj Nair, Avani Narayan, Deepak Narayanan, Benjamin Newman, Allen Nie, Juan Carlos Niebles, Hamed Nilforoshan, J. F. Nyarko, Giray Ogut, Laurel Orr, Isabel Papadimitriou, Joon Sung Park, Chris Piech, Eva Portelance, Christopher Potts, Aditi Raghunathan, Robert Reich, Hongyu Ren, Frieda Rong, Yusuf H. Roohani, Camilo Ruiz, Jack Ryan, Christopher R’e, Dorsa Sadigh, Shiori Sagawa, Keshav Santhanam, Andy Shih, Krishna Parasuram Srinivasan, Alex Tamkin, Rohan Taori, Armin W. Thomas, Florian Tramèr, Rose E. Wang, William Wang, Bohan Wu, Jiajun Wu, Yuhuai Wu, Sang Michael Xie, Michihiro Yasunaga, Jiaxuan You, Matei A. Zaharia, Michael Zhang, Tianyi Zhang, Xikun Zhang, Yuhui Zhang, Lucia Zheng, Kaitlyn Zhou, and Percy Liang. 2021. On the opportunities and risks of foundation models. *arXiv*, arXiv:2108.07258.
- Bowman, Samuel. 2022. The dangers of underclaiming: Reasons for caution when

- reporting how NLP systems fail. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 7484–7499.
- Brandl, Stephanie, Ruixiang Cui, and Anders Søgaard. 2022. How conservative are language models? adapting to the introduction of gender-neutral pronouns. In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 3624–3630.
- Brown, Hannah, Katherine Lee, Fatemehsadat Mireshghallah, Reza Shokri, and Florian Tramèr. 2022. What does it mean for a language model to preserve privacy? In *Proceedings of the ACM Conference on Fairness, Accountability, and Transparency*.
- Brown, Tom, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel Ziegler, Jeffrey Wu, Clemens Winter, Chris Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam McCandlish, Alec Radford, Ilya Sutskever, and Dario Amodei. 2020. Language models are few-shot learners. In *Advances in Neural Information Processing Systems*, volume 33, pages 1877–1901.
- Broyde, Joshua and Claire Palmer. 2021. Build a medical sentence matching application using BERT and Amazon SageMaker. *AWS Machine Learning Blog*.
- Cai, Xingyu, Jiaji Huang, Yu-Lan Bian, and Kenneth Ward Church. 2021. Isotropy in the contextual embedding space: Clusters and manifolds. In *International Conference on Learning Representations*.
- Cao, Boxi, Hongyu Lin, Xianpei Han, Fangchao Liu, and Le Sun. 2022. Can prompt probe pretrained language models? understanding the invisible risks from a causal view. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 5796–5808.
- Cao, Boxi, Hongyu Lin, Xianpei Han, Le Sun, Lingyong Yan, Meng Liao, Tong Xue, and Jin Xu. 2021. Knowledgeable or educated guess? revisiting language models as knowledge bases. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 1860–1874.
- Carlini, Nicholas, Daphne Ippolito, Matthew Jagielski, Katherine Lee, Florian Tramèr, and Chiyuan Zhang. 2023. Quantifying memorization across neural language models. In *International Conference on Learning Representations*.
- Carlini, Nicholas, Florian Tramèr, Eric Wallace, Matthew Jagielski, Ariel Herbert-Voss, Katherine Lee, Adam Roberts, Tom Brown, Dawn Song, Ulfar Erlingsson, Alina Oprea, and Colin Raffel. 2021. Extracting training data from large language models. In *USENIX Security Symposium*, pages 2633–2650.
- Carnie, Andrew. 2002. *Syntax: A Generative Introduction*. Blackwell.
- Caron, Graham and Shashank Srivastava. 2022. Identifying and manipulating the personality traits of language models. *arXiv*, arXiv:2212.10276.
- Chang, Tyler, Zhuowen Tu, and Benjamin Bergen. 2022. The geometry of multilingual language model representations. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 119–136.
- Chang, Tyler, Yifan Xu, Weijian Xu, and Zhuowen Tu. 2021. Convolutions and self-attention: Re-interpreting relative positions in pre-trained language models. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 4322–4333.
- Chang, Tyler A. and Benjamin K. Bergen. 2022. Word acquisition in neural language models. *Transactions of the Association for Computational Linguistics*, 10:1–16.
- Chaves, Rui P. and Stephanie N. Richter. 2021. Look at that! BERT can be easily distracted from paying attention to morphosyntax. In *Proceedings of the Society for Computation in Linguistics 2021*, pages 28–38.
- Chen, Kaiping, Anqi Shao, Jirayu Burapachee, and Yixuan Li. 2022. A critical appraisal of equity in conversational AI: Evidence from auditing GPT-3’s dialogues with different publics on climate change and Black Lives Matter. *ArXiv*, arXiv:2209.13627.
- Chen, Mark, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde, Jared Kaplan, Harrison Edwards, Yura Burda, Nicholas Joseph, Greg Brockman, Alex

- Ray, Raul Puri, Gretchen Krueger, Michael Petrov, Heidy Khlaaf, Girish Sastry, Pamela Mishkin, Brooke Chan, Scott Gray, Nick Ryder, Mikhail Pavlov, Alethea Power, Lukasz Kaiser, Mohammad Bavarian, Clemens Winter, Philippe Tillet, F. Such, D. Cummings, Matthias Plappert, Fotios Chantzis, Elizabeth Barnes, Ariel Herbert-Voss, William H. Guss, Alex Nichol, I. Babuschkin, S. Balaji, Shantanu Jain, A. Carr, J. Leike, Joshua Achiam, Vedant Misra, Evan Morikawa, Alec Radford, M. Knight, Miles Brundage, Mira Murati, Katie Mayer, P. Welinder, Bob McGrew, Dario Amodei, Sam McCandlish, Ilya Sutskever, and Wojciech Zaremba. 2021. Evaluating large language models trained on code. *ArXiv*, arXiv:2107.03374.
- Chiang, Cheng-Han, Sung-Feng Huang, and Hung-yi Lee. 2020. Pretrained language model embryology: The birth of ALBERT. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 6813–6828.
- Chiang, Ting-Rui and Yun-Nung Chen. 2021. Relating neural text degeneration to exposure bias. In *Proceedings of the Fourth BlackboxNLP Workshop on Analyzing and Interpreting Neural Networks for NLP*, pages 228–239.
- Cho, Won Ik, Emmanuele Chersoni, Yu-Yin Hsu, and Chu-Ren Huang. 2021. Modeling the influence of verb aspect on the activation of typical event locations with BERT. In *Findings of the Association for Computational Linguistics: ACL-IJCNLP 2021*, pages 2922–2929.
- Choenni, Rochelle, Ekaterina Shutova, and Robert van Rooij. 2021. Stepmothers are mean and academics are pretentious: What do pretrained language models learn about you? In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 1477–1491.
- Choshen, Leshem, Guy Hacohen, Daphna Weinshall, and Omri Abend. 2022. The grammar-learning trajectories of neural language models. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 8281–8297.
- Choudhury, Monojit and Amit Deshpande. 2021. How linguistically fair are multilingual pre-trained language models? In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 35, pages 12710–12718.
- Chowdhery, Aakanksha, Sharan Narang, Jacob Devlin, Maarten Bosma, Gaurav Mishra, Adam Roberts, Paul Barham, Hyung Won Chung, Charles Sutton, Sebastian Gehrmann, Parker Schuh, Kensen Shi, Sasha Tsvyashchenko, Joshua Maynez, Abhishek Rao, Parker Barnes, Yi Tay, Noam M. Shazeer, Vinodkumar Prabhakaran, Emily Reif, Nan Du, Benton C. Hutchinson, Reiner Pope, James Bradbury, Jacob Austin, Michael Isard, Guy Gur-Ari, Pengcheng Yin, Toju Duke, Anselm Levskaya, Sanjay Ghemawat, Sunipa Dev, Henryk Michalewski, Xavier Garcia, Vedant Misra, Kevin Robinson, Liam Fedus, Denny Zhou, Daphne Ippolito, David Luan, Hyeontaek Lim, Barret Zoph, Alexander Spiridonov, Ryan Sepassi, David Dohan, Shivani Agrawal, Mark Omernick, Andrew M. Dai, Thanumalayan Sankaranarayanan Pillai, Marie Pellat, Aitor Lewkowycz, Erica Moreira, Rewon Child, Oleksandr Polozov, Katherine Lee, Zongwei Zhou, Xuezhi Wang, Brennan Saeta, Mark Díaz, Orhan Firat, Michele Catasta, Jason Wei, Kathleen S. Meier-Hellstern, Douglas Eck, Jeff Dean, Slav Petrov, and Noah Fiedel. 2022. PaLM: Scaling language modeling with Pathways. *arXiv*, arXiv:2204.02311.
- Chuang, Chengyu and Yi Yang. 2022. Buy tesla, sell ford: Assessing implicit stock market preference in pre-trained language models. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, pages 100–105.
- Cifka, Ondřej and Antoine Liutkus. 2022. Black-box language model explanation by context length probing. *ArXiv*, arXiv:2212.14815.
- Clark, Elizabeth, Tal August, Sofia Serrano, Nikita Haduong, Suchin Gururangan, and Noah A. Smith. 2021. All that’s ‘human’ is not gold: Evaluating human evaluation of generated text. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 7282–7296.
- Clark, Kevin, Urvashi Khandelwal, Omer Levy, and Christopher D. Manning. 2019. What does BERT look at? an analysis of BERT’s attention. In *Proceedings of the 2019 ACL Workshop BlackboxNLP: Analyzing and Interpreting Neural Networks for NLP*, pages 276–286.
- Cohen, Aaron Daniel, Adam Roberts, Alejandra Molina, Alena Butryna, Alicia Jin, Apoorv Kulshreshtha, Ben

- Hutchinson, Ben Zevenbergen, Blaise Hilary Aguera-Arcas, Chung ching Chang, Claire Cui, Cosmo Du, Daniel De Freitas Adiwardana, Dehao Chen, Dmitry (Dima) Lepikhin, Ed H. Chi, Erin Hoffman-John, Heng-Tze Cheng, Hongrae Lee, Igor Krivokon, James Qin, Jamie Hall, Joe Fenton, Johnny Soraker, Kathy Meier-Hellstern, Kristen Olson, Lora Moiso Aroyo, Maarten Paul Bosma, Marc Joseph Pickett, Marcelo Amorim Menegali, Marian Croak, Mark Díaz, Matthew Lamm, Maxim Krikun, Meredith Ringel Morris, Noam Shazeer, Quoc V. Le, Rachel Bernstein, Ravi Rajakumar, Ray Kurzweil, Romal Thoppilan, Steven Zheng, Taylor Bos, Toju Duke, Tulsee Doshi, Vincent Y. Zhao, Vinodkumar Prabhakaran, Will Rusch, YaGuang Li, Yanping Huang, Yanqi Zhou, Yuanzhong Xu, and Zhifeng Chen. 2022. LaMDA: Language models for dialog applications. *ArXiv*, arXiv:2201.08239.
- Comşa, Iulia, Julian Eisenschlos, and Srini Narayanan. 2022. MiQA: A benchmark for inference on metaphorical questions. In *Proceedings of the 2nd Conference of the Asia-Pacific Chapter of the Association for Computational Linguistics and the 12th International Joint Conference on Natural Language Processing (Volume 2: Short Papers)*, pages 373–381.
- Cong, Yan. 2022. Psycholinguistic diagnosis of language models' commonsense reasoning. In *Proceedings of the First Workshop on Commonsense Representation and Reasoning (CSRR 2022)*, pages 17–22.
- Czarnowska, Paula, Yogarshi Vyas, and Kashif Shah. 2021. Quantifying social biases in NLP: A generalization and empirical comparison of extrinsic fairness metrics. *Transactions of the Association for Computational Linguistics*, 9:1249–1267.
- Czinczoll, Tamara, Helen Yannakoudakis, Pushkar Mishra, and Ekaterina Shutova. 2022. Scientific and creative analogies in pretrained language models. In *Findings of the Association for Computational Linguistics: EMNLP 2022*, pages 2094–2100.
- Dąbrowska, Ewa. 2015. What exactly is Universal Grammar, and has anyone seen it? *Frontiers in Psychology*, 6.
- Dai, Damai, Yutao Sun, Li Dong, Yaru Hao, Zhifang Sui, and Furu Wei. 2022. Why can GPT learn in-context? language models secretly perform gradient descent as meta-optimizers. *ArXiv*, arXiv:2212.10559.
- Dai, Zihang, Zhilin Yang, Yiming Yang, Jaime Carbonell, Quoc Le, and Ruslan Salakhutdinov. 2019. Transformer-XL: Attentive language models beyond a fixed-length context. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 2978–2988.
- Dar, Guy, Mor Geva, Ankit Gupta, and Jonathan Berant. 2022. Analyzing Transformers in embedding space. *ArXiv*, arXiv:2209.02535.
- Dasgupta, Ishita, Andrew Lampinen, Stephanie Chan, Antonia Creswell, Dharshan Kumaran, James McClelland, and Felix Hill. 2022. Language models show human-like content effects on reasoning. *arXiv*, arXiv:2207.07051.
- Davis, Forrest and Marten van Schijndel. 2020. Discourse structure interacts with reference but not syntax in neural language models. In *Proceedings of the 24th Conference on Computational Natural Language Learning*, pages 396–407.
- Davison, Joe, Joshua Feldman, and Alexander Rush. 2019. Commonsense knowledge mining from pretrained models. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 1173–1178.
- De Bruyn, Maxime, Ehsan Lotfi, Jeska Buhmann, and Walter Daelemans. 2022. Is it smaller than a tennis ball? language models play the game of twenty questions. In *Proceedings of the Fifth BlackboxNLP Workshop on Analyzing and Interpreting Neural Networks for NLP*, pages 80–90.
- Dev, Sunipa, Masoud Monajatipoor, Anaelia Ovalle, Arjun Subramonian, Jeff Phillips, and Kai-Wei Chang. 2021. Harms of gender exclusivity and challenges in non-binary representation in language technologies. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 1968–1994.
- Dev, Sunipa, Emily Sheng, Jieyu Zhao, Aubrie Amstutz, Jiao Sun, Yu Hou, Mattie Sanseverino, Jiin Kim, Akihiro Nishi, Nanyun Peng, and Kai-Wei Chang. 2022. On measures of biases and harms in NLP. In *Findings of the Association for Computational Linguistics: ACL-IJCNLP 2022*, pages 246–267.
- Devlin, Jacob, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2019. BERT: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for*

- Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 4171–4186.
- Dhamala, Jwala, Tony Sun, Varun Kumar, Satyapriya Krishna, Yada Pruksachatkun, Kai-Wei Chang, and Rahul Gupta. 2021. BOLD: Dataset and metrics for measuring biases in open-ended language generation. In *Proceedings of the ACM Conference on Fairness, Accountability, and Transparency*, page 862–872, Association for Computing Machinery, New York, NY, USA.
- Du, Mengnan, Fengxiang He, Na Zou, Dacheng Tao, and Xia Hu. 2022a. Shortcut learning of large language models in natural language understanding: A survey. *ArXiv*, arXiv:2208.11857.
- Du, Yifan, Zikang Liu, Junyi Li, and Wayne Xin Zhao. 2022b. A survey of vision-language pre-trained models. In *Proceedings of the International Joint Conference on Artificial Intelligence*, pages 5436–5443. Survey Track.
- Dufter, Philipp, Martin Schmitt, and Hinrich Schütze. 2022. Position information in transformers: An overview. *Computational Linguistics*, 48(3):733–763.
- Dugan, Liam, Daphne Ippolito, Arun Kirubarajan, Sherry Shi, and Chris Callison-Burch. 2023. Real or fake text?: Investigating human ability to detect boundaries between human-written and machine-generated text. In *Proceedings of the AAAI Conference on Artificial Intelligence*.
- Elazar, Yanai, Nora Kassner, Shauli Ravfogel, Amir Feder, Abhilasha Ravichander, Marius Mosbach, Yonatan Belinkov, Hinrich Schütze, and Yoav Goldberg. 2022. Measuring causal effects of data statistics on language model’s ‘factual’ predictions. *ArXiv*, arXiv:2207.14251.
- Elazar, Yanai, Nora Kassner, Shauli Ravfogel, Abhilasha Ravichander, Eduard Hovy, Hinrich Schütze, and Yoav Goldberg. 2021. Measuring and improving consistency in pretrained language models. *Transactions of the Association for Computational Linguistics*, 9:1012–1031.
- Ettinger, Allyson. 2020. What BERT is not: Lessons from a new suite of psycholinguistic diagnostics for language models. *Transactions of the Association for Computational Linguistics*, 8:34–48.
- Fedus, William, Barret Zoph, and Noam Shazeer. 2022. Switch Transformers: Scaling to trillion parameter models with simple and efficient sparsity. *Journal of Machine Learning Research*, 23.
- Felkner, Virginia K., Ho-Chun Herbert Chang, Eugene Jang, and Jonathan May. 2022. Towards WinoQueer: Developing a benchmark for anti-queer bias in large language models. In *Queer in AI Workshop*.
- Finlayson, Matthew, Aaron Mueller, Sebastian Gehrmann, Stuart Shieber, Tal Linzen, and Yonatan Belinkov. 2021. Causal analysis of syntactic agreement mechanisms in neural language models. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 1828–1843.
- Frank, Michael and Noah Goodman. 2012. Predicting pragmatic reasoning in language games. *Science*, 336(6084).
- Freitag, Markus and Yaser Al-Onaizan. 2017. Beam search strategies for neural machine translation. In *Proceedings of the First Workshop on Neural Machine Translation*, pages 56–60.
- Fricke, Suzanne. 2018. Semantic Scholar. *Journal of the Medical Library Association*, 106(1).
- Fujisawa, Ippei and Ryota Kanai. 2022. Logical tasks for measuring extrapolation and rule comprehension. *ArXiv*, arXiv:2211.07727.
- Ganguli, Deep, Danny Hernandez, Liane Lovitt, Amanda Askell, Yuntao Bai, Anna Chen, Tom Conerly, Nova Dassarma, Dawn Drain, Nelson Elhage, Sheer El Showk, Stanislav Fort, Zac Hatfield-Dodds, Tom Henighan, Scott Johnston, Andy Jones, Nicholas Joseph, Jackson Kernian, Shauna Kravec, Ben Mann, Neel Nanda, Kamal Ndousse, Catherine Olsson, Daniela Amodei, Tom Brown, Jared Kaplan, Sam McCandlish, Christopher Olah, Dario Amodei, and Jack Clark. 2022a. Predictability and surprise in large generative models. In *Proceedings of the ACM Conference on Fairness, Accountability, and Transparency*, page 1747–1764, Association for Computing Machinery, New York, NY, USA.
- Ganguli, Deep, Liane Lovitt, John Kernion, Amanda Askell, Yuntao Bai, Saurav Kadavath, Benjamin Mann, Ethan Perez, Nicholas Schiefer, Kamal Ndousse, Andy Jones, Sam Bowman, Anna Chen, Tom Conerly, Nova DasSarma, Dawn Drain, Nelson Elhage, Sheer El-Showk, Stanislav Fort, Zachary Dodds, T. J. Henighan, Danny Hernandez, Tristan Hume, Josh Jacobson, Scott Johnston, Shauna Kravec, Catherine Olsson, Sam Ringer, Eli

- Tran-Johnson, Dario Amodei, Tom B. Brown, Nicholas Joseph, Sam McCandlish, Christopher Olah, Jared Kaplan, and Jack Clark. 2022b. Red teaming language models to reduce harms: Methods, scaling behaviors, and lessons learned. *ArXiv*, arXiv:2209.07858.
- Gauthier, Jon, Jennifer Hu, Ethan Wilcox, Peng Qian, and Roger Levy. 2020. SyntaxGym: An online platform for targeted evaluation of language models. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics: System Demonstrations*, pages 70–76.
- Geeraerts, Dirk. 2017. Lexical semantics. *Oxford Research Encyclopedia of Linguistics*.
- Gehman, Samuel, Suchin Gururangan, Maarten Sap, Yejin Choi, and Noah A. Smith. 2020. RealToxicityPrompts: Evaluating neural toxic degeneration in language models. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, pages 3356–3369.
- Geiger, Atticus, Kyle Richardson, and Christopher Potts. 2020. Neural natural language inference models partially embed theories of lexical entailment and negation. In *Proceedings of the Third BlackboxNLP Workshop on Analyzing and Interpreting Neural Networks for NLP*, pages 163–173.
- Geva, Mor, Avi Caciularu, Kevin Wang, and Yoav Goldberg. 2022. Transformer feed-forward layers build predictions by promoting concepts in the vocabulary space. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 30–45.
- Geva, Mor, Roei Schuster, Jonathan Berant, and Omer Levy. 2021. Transformer feed-forward layers are key-value memories. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 5484–5495.
- Goldberg, Yoav. 2019. Assessing BERT’s syntactic abilities. *ArXiv*, arXiv:1901.05287.
- Grice, H. P. 1975. Logic and conversation. *Syntax and Semantics: Vol. 3: Speech Acts*.
- Griciūtė, Bernadeta, Marc Tanti, and Lucia Donatelli. 2022. On the cusp of comprehensibility: Can language models distinguish between metaphors and nonsense? In *Proceedings of the 3rd Workshop on Figurative Language Processing (FLP)*, pages 173–177.
- Groenwold, Sophie, Lily Ou, Aesha Parekh, Samhita Honnavalli, Sharon Levy, Diba Mirza, and William Yang Wang. 2020. Investigating African-American Vernacular English in transformer-based text generation. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 5877–5883.
- Gubelmann, Reto and Siegfried Handschuh. 2022. Context matters: A pragmatic study of PLMs’ negation understanding. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 4602–4621.
- Hagendorff, Thilo, Sarah Fabi, and Michal Kosinski. 2022. Machine intuition: Uncovering human-like intuitive decision-making in GPT-3.5. *ArXiv*, arXiv:2212.05206.
- Hahn, Michael. 2020. Theoretical limitations of self-attention in neural sequence models. *Transactions of the Association for Computational Linguistics*, 8:156–171.
- Han, Simeng, Hailey Schoelkopf, Yilun Zhao, Zhenting Qi, Martin Riddell, Luke Benson, Lucy Sun, Ekaterina Zubova, Yujie Qiao, Matthew Burtell, David Peng, Jonathan Fan, Yixin Liu, Brian Wong, Malcolm Sailor, Ansong Ni, Linyong Nan, Jungo Kasai, Tao Yu, Rui Zhang, Shafiq Joty, Alexander R. Fabbri, Wojciech Kryscinski, Xi Victoria Lin, Caiming Xiong, and Dragomir Radev. 2022. FOLIO: Natural language reasoning with first-order logic. *ArXiv*, arXiv:2209.00840.
- Hanna, Michael and David Mareček. 2021. Analyzing BERT’s knowledge of hypernymy via prompting. In *Proceedings of the Fourth BlackboxNLP Workshop on Analyzing and Interpreting Neural Networks for NLP*, pages 275–282.
- Hassan, Saad, Matt Huenerfauth, and Cecilia Ovesdotter Alm. 2021. Unpacking the interdependent systems of discrimination: Ableist bias in NLP systems through an intersectional lens. In *Findings of the Association for Computational Linguistics: EMNLP 2021*, pages 3116–3123.
- Haviv, Adi, Ori Ram, Ofir Press, Peter Izsak, and Omer Levy. 2022. Transformer language models without positional encodings still learn positional information. In *Findings of the Association for Computational Linguistics: EMNLP 2022*, pages 1382–1390.
- Hawkins, Robert, Takateru Yamakoshi, Thomas Griffiths, and Adele Goldberg. 2020. Investigating representations of verb bias in neural language models. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 4653–4663.

- He, Qianyu, Sijie Cheng, Zhixu Li, Rui Xie, and Yanghua Xiao. 2022a. Can pre-trained language models interpret similes as smart as human? In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 7875–7887.
- He, Xuanli, Qionghai Xu, Lingjuan Lyu, Fangzhao Wu, and Chenguang Wang. 2022b. Protecting intellectual property of language generation APIs with lexical watermark. In *Proceedings of the AAAI Conference on Artificial Intelligence*.
- Heidenreich, Hunter Scott and Jake Ryland Williams. 2021. The Earth is flat and the Sun is not a star: The susceptibility of GPT-2 to universal adversarial triggers. In *Proceedings of the 2021 AAAI/ACM Conference on AI, Ethics, and Society*, page 566–573, Association for Computing Machinery, New York, NY, USA.
- Hendrycks, Dan, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. 2021a. Measuring massive multitask language understanding. In *International Conference on Learning Representations*.
- Hendrycks, Dan, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, and Jacob Steinhardt. 2021b. Measuring mathematical problem solving with the MATH dataset. In *Advances in Neural Information Processing Systems*.
- Hernandez, Danny, Tom B. Brown, Tom Conerly, Nova DasSarma, Dawn Drain, Sheer El-Showk, Nelson Elhage, Zac Hatfield-Dodds, Tom Henighan, Tristan Hume, Scott Johnston, Benjamin Mann, Christopher Olah, Catherine Olsson, Dario Amodei, Nicholas Joseph, Jared Kaplan, and Sam McCandlish. 2022. Scaling laws and interpretability of learning from repeated data. *ArXiv*, arXiv:2205.10487.
- Hershcovich, Daniel, Stella Frank, Heather Lent, Miryam de Lhoneux, Mostafa Abdou, Stephanie Brandl, Emanuele Bugliarello, Laura Cabello Piqueras, Ilias Chalkidis, Ruixiang Cui, Constanza Fierro, Katerina Margatina, Phillip Rust, and Anders Søgaard. 2022. Challenges and strategies in cross-cultural NLP. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 6997–7013.
- Hoffmann, Jordan, Sebastian Borgeaud, Arthur Mensch, Elena Buchatskaya, Trevor Cai, Eliza Rutherford, Diego de las Casas, Lisa Anne Hendricks, Johannes Welbl, Aidan Clark, Tom Hennigan, Eric Noland, Katherine Millican, George van den Driessche, Bogdan Damoc, Aurelia Guy, Simon Osindero, Karen Simonyan, Erich Elsen, Oriol Vinyals, Jack William Rae, and Laurent Sifre. 2022. Training compute-optimal large language models. In *Advances in Neural Information Processing Systems*.
- Holtzman, Ari, Jan Buys, Li Du, Maxwell Forbes, and Yejin Choi. 2020. The curious case of neural text degeneration. In *International Conference on Learning Representations*.
- Hossain, Md Mosharaf, Dhivya Chinnappa, and Eduardo Blanco. 2022. An analysis of negation in natural language understanding corpora. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, pages 716–723.
- Hossain, Md Mosharaf, Venelin Kovatchev, Pranoy Dutta, Tiffany Kao, Elizabeth Wei, and Eduardo Blanco. 2020. An analysis of natural language inference benchmarks through the lens of negation. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 9106–9118.
- Hosseini, Arian, Ankit Vani, Dzmitry Bahdanau, Alessandro Sordani, and Aaron Courville. 2022. On the compositional generalization gap of in-context learning. In *Proceedings of the Fifth BlackboxNLP Workshop on Analyzing and Interpreting Neural Networks for NLP*, pages 272–280.
- Hu, Jennifer, Sherry Yong Chen, and Roger Levy. 2020. A closer look at the performance of neural language models on reflexive anaphor licensing. In *Proceedings of the Society for Computation in Linguistics 2020*, pages 323–333.
- Hu, Jennifer, Sammy Floyd, Olessia Jouravlev, Evelina Fedorenko, and Edward Gibson. 2022. A fine-grained comparison of pragmatic language understanding in humans and language models. *ArXiv*, arXiv:2212.06801.
- Hu, Jennifer, Jon Gauthier, Peng Qian, Ethan Wilcox, and Roger Levy. 2020. A systematic assessment of syntactic generalization in neural language models. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 1725–1744.
- Huang, Jie, Hanyin Shao, and Kevin Chen-Chuan Chang. 2022. Are large pre-trained language models leaking your personal information? In *Findings of the*

- Association for Computational Linguistics: EMNLP 2022*, pages 2038–2047.
- Huebner, Philip A., Elior Sulem, Fisher Cynthia, and Dan Roth. 2021. BabyBERTa: Learning more grammar with small-scale child-directed language. In *Proceedings of the 25th Conference on Computational Natural Language Learning*, pages 624–646.
- Hupkes, Dieuwke, Mario Giulianelli, Verna Dankers, Mikel Artetxe, Yanai Elazar, Tiago Pimentel, Christos Christodoulopoulos, Karim Lasri, Naomi Saphra, Arabella Sinclair, Dennis Ulmer, Florian Schottnmann, Khuyagbaatar Batsuren, Kaiser Sun, Koustuv Sinha, Leila Khalatbari, Maria Ryskina, Rita Frieske, Ryan Cotterell, and Zhijing Jin. 2022. State-of-the-art generalisation research in NLP: A taxonomy and review. *arXiv*, arXiv:2210.03050.
- Huynh, Hien, Tomas O. Lentz, and Emiel van Miltenburg. 2022. Implicit causality in GPT-2: A case study. *ArXiv*, arXiv:2212.04348.
- Ippolito, Daphne, Daniel Duckworth, Chris Callison-Burch, and Douglas Eck. 2020. Automatic detection of generated text is easiest when humans are fooled. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 1808–1822.
- Ippolito, Daphne, Florian Tramèr, Milad Nasr, Chiyuan Zhang, Matthew Jagielski, Katherine Lee, Christopher A. Choquette-Choo, and Nicholas Carlini. 2022. Preventing verbatim memorization in language models gives a false sense of privacy. *ArXiv*, arXiv:2210.17546.
- Iyer, Srinivas, Xiaojuan Lin, Ramakanth Pasunuru, Todor Mihaylov, Daniel Simig, Ping Yu, Kurt Shuster, Tianlu Wang, Qing Liu, Punit Singh Koura, Xian Li, Brian O'Horo, Gabriel Pereyra, Jeff Wang, Christopher Dewan, Asli Celikyilmaz, Luke Zettlemoyer, and Veselin Stoyanov. 2022. OPT-IML: Scaling language model instruction meta learning through the lens of generalization. *ArXiv*, arXiv:2212.12017.
- Jain, Sarthak and Byron C. Wallace. 2019. Attention is not Explanation. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 3543–3556.
- Jakesch, Maurice, Jeffrey T. Hancock, and Mor Naaman. 2023. Human heuristics for AI-generated language are flawed. *Proceedings of the National Academy of Sciences*, 120(11).
- Jang, Joel, Seonghyeon Ye, and Minjoon Seo. 2022. Can large language models truly understand prompts? a case study with negated prompts. In *Workshop on Transfer Learning for Natural Language Processing*.
- Jawahar, Ganesh, Muhammad Abdul-Mageed, and Laks Lakshmanan, V.S. 2020. Automatic detection of machine generated text: A critical survey. In *Proceedings of the 28th International Conference on Computational Linguistics*, pages 2296–2309, International Committee on Computational Linguistics, Barcelona, Spain (Online).
- Jiang, Chengyue, Zhonglin Nian, Kaihao Guo, Shanbo Chu, Yingong Zhao, Libin Shen, and Kewei Tu. 2020a. Learning numeral embedding. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, pages 2586–2599.
- Jiang, Guangyuan, Manjie Xu, Song-Chun Zhu, Wenjuan Han, Chi Zhang, and Yixin Zhu. 2022. MPI: Evaluating and inducing personality in pre-trained language models. *ArXiv*, arXiv:2206.07550.
- Jiang, Tianyu and Ellen Riloff. 2021. Learning prototypical functions for physical artifacts. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 6941–6951.
- Jiang, Zhengbao, Frank F. Xu, Jun Araki, and Graham Neubig. 2020b. How can we know what language models know? *Transactions of the Association for Computational Linguistics*, 8:423–438.
- Jin, Zhijing, Sydney Levine, Fernando Gonzalez Adauro, Ojasv Kamal, Maarten Sap, Mrinmaya Sachan, Rada Mihalcea, Joshua B. Tenenbaum, and Bernhard Schölkopf. 2022a. When to make exceptions: Exploring language models as accounts of human moral judgment. In *Advances in Neural Information Processing Systems*.
- Jin, Zijia, Xingyu Zhang, Mo Yu, and Lifu Huang. 2022b. Probing script knowledge from pre-trained models. In *Proceedings of the Workshop on Unimodal and Multimodal Induction of Linguistic Structures (UM-IoS)*, pages 87–93.
- Johnson, Rebecca L, Giada Pistilli, Natalia Menéndez-González, Leslye Denisse Dias Duran, Enrico Panai, Julija Kalpokiene, and Donald Jay Bertulfo. 2022. The ghost in the machine has an American accent: Value conflict in GPT-3. *ArXiv*,

- arXiv:2203.07785.
- Johnson, Steven. 2022. A.I. is mastering language. Should we trust what it says? *The New York Times*.
- Jones, Cameron and Benjamin Bergen. 2021. The role of physical inference in pronoun resolution. In *Proceedings of the Annual Meeting of the Cognitive Science Society*, volume 43.
- Jones, Cameron R., Tyler A. Chang, Seana Coulson, James Michaelov, Sean Trott, and Benjamin Bergen. 2022. Distributional semantics still can't account for affordances. In *Proceedings of the Annual Meeting of the Cognitive Science Society*, volume 44.
- Joshi, Pratik, Sebastin Santy, Amar Budhiraja, Kalika Bali, and Monojit Choudhury. 2020. The state and fate of linguistic diversity and inclusion in the NLP world. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 6282–6293.
- Kabbara, Jad and Jackie Chi Kit Cheung. 2022. Investigating the performance of transformer-based NLI models on presuppositional inferences. In *Proceedings of the 29th International Conference on Computational Linguistics*, pages 779–785, International Committee on Computational Linguistics, Gyeongju, Republic of Korea.
- Kadavath, Saurav, Tom Conerly, Amanda Askell, Tom Henighan, Dawn Drain, Ethan Perez, Nicholas Schiefer, Zachary Dodds, Nova DasSarma, Eli Tran-Johnson, Scott Johnston, Sheer El-Showk, Andy Jones, Nelson Elhage, Tristan Hume, Anna Chen, Yuntao Bai, Sam Bowman, Stanislaw Fort, Deep Ganguli, Danny Hernandez, Josh Jacobson, John Kernion, Shauna Kravec, Liane Lovitt, Kamal Ndousse, Catherine Olsson, Sam Ringer, Dario Amodei, Tom B. Brown, Jack Clark, Nicholas Joseph, Benjamin Mann, Sam McCandlish, Christopher Olah, and Jared Kaplan. 2022. Language models (mostly) know what they know. *ArXiv*, arXiv:2207.05221.
- Kalo, Jan-Christoph and Leandra Fichtel. 2022. KAMEL: Knowledge analysis with multitoken entities in language models. In *4th Conference on Automated Knowledge Base Construction*.
- Kandpal, Nikhil, Haikang Deng, Adam Roberts, Eric Wallace, and Colin Raffel. 2022. Large language models struggle to learn long-tail knowledge. *ArXiv*, arXiv:2211.08411.
- Kandpal, Nikhil, Eric Wallace, and Colin Raffel. 2022. Deduplicating training data mitigates privacy risks in language models. In *International Conference on Machine Learning*.
- Kaplan, Jared, Sam McCandlish, Tom Henighan, Tom B. Brown, Benjamin Chess, Rewon Child, Scott Gray, Alec Radford, Jeff Wu, and Dario Amodei. 2020. Scaling laws for neural language models. *ArXiv*, arXiv:2001.08361.
- Kassner, Nora, Benno Krojer, and Hinrich Schütze. 2020. Are pretrained language models symbolic reasoners over knowledge? In *Proceedings of the 24th Conference on Computational Natural Language Learning*, pages 552–564.
- Kassner, Nora and Hinrich Schütze. 2020. Negated and misprimed probes for pretrained language models: Birds can talk, but cannot fly. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 7811–7818.
- Katz, Uri, Mor Geva, and Jonathan Berant. 2022. Inferring implicit relations in complex questions with language models. In *Findings of the Association for Computational Linguistics: EMNLP 2022*, pages 2548–2566.
- Kauf, Carina, Anna A. Ivanova, Giulia Rambelli, Emmanuele Chersoni, Jingyuan S. She, Zawad Chowdhury, Evelina Fedorenko, and Alessandro Lenci. 2022. Event knowledge in large language models: The gap between the impossible and the unlikely. *ArXiv*, arXiv:2212.01488.
- Kavumba, Pride, Naoya Inoue, Benjamin Heinzerling, Keshav Singh, Paul Reiser, and Kentarou Inui. 2020. Balanced COPA: Countering superficial cues in causal reasoning. *Association for Natural Language Processing*.
- Kavumba, Pride, Ryo Takahashi, and Yusuke Oda. 2022. Are prompt-based models clueless? In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 2333–2352.
- Kementchedjhieva, Yova, Mark Anderson, and Anders Søgaard. 2021. John praised Mary because _he_? implicit causality bias and its interaction with explicit cues in LMs. In *Findings of the Association for Computational Linguistics: ACL-IJCNLP 2021*, pages 4859–4871.
- Kharitonov, Eugene, Marco Baroni, and Dieuwke Hupkes. 2021. How BPE affects memorization in Transformers. *ArXiv*, arXiv:2110.02782.

- Kim, Sanghee J., Lang Yu, and Allyson Ettinger. 2022. "no, they did not": Dialogue response dynamics in pre-trained language models. In *Proceedings of the 29th International Conference on Computational Linguistics*, pages 863–874, International Committee on Computational Linguistics, Gyeongju, Republic of Korea.
- Kirchenbauer, John, Jonas Geiping, Yuxin Wen, Jonathan Katz, Ian Miers, and Tom Goldstein. 2023. A watermark for large language models. *ArXiv*, arXiv:2301.10226.
- Kirk, Hannah Rose, Yennie Jun, Filippo Volpin, Haider Iqbal, Elias Benussi, Frederic Dreyer, Aleksandar Shtedritski, and Yuki Asano. 2021. Bias out-of-the-box: An empirical analysis of intersectional occupational biases in popular generative language models. In *Advances in Neural Information Processing Systems*, volume 34, pages 2611–2624.
- Ko, Wei-Jen and Junyi Jessy Li. 2020. Assessing discourse relations in language generation from GPT-2. In *Proceedings of the 13th International Conference on Natural Language Generation*, pages 52–59.
- Kojima, Takeshi, Shixiang Shane Gu, Machel Reid, Yutaka Matsuo, and Yusuke Iwasawa. 2022. Large language models are zero-shot reasoners. In *Advances in Neural Information Processing Systems*.
- Kovaleva, Olga, Alexey Romanov, Anna Rogers, and Anna Rumshisky. 2019. Revealing the dark secrets of BERT. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 4365–4374.
- Kudo, Taku. 2018. Subword regularization: Improving neural network translation models with multiple subword candidates. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 66–75.
- Kurita, Keita, Nidhi Vyas, Ayush Pareek, Alan W Black, and Yulia Tsvetkov. 2019. Measuring bias in contextualized word representations. In *Proceedings of the First Workshop on Gender Bias in Natural Language Processing*, pages 166–172.
- Kwon, Sunjae, Cheongwoong Kang, Jiyeon Han, and Jaesik Choi. 2019. Why do masked neural language models still need common sense knowledge? *ArXiv*, arXiv:1911.03024.
- Lakhotia, Kushal, Eugene Kharitonov, Wei-Ning Hsu, Yossi Adi, Adam Polyak, Benjamin Bolte, Tu-Anh Nguyen, Jade Copet, Alexei Baevski, Abdelrahman Mohamed, and Emmanuel Dupoux. 2021. On generative spoken language modeling from raw audio. *Transactions of the Association for Computational Linguistics*, 9:1336–1354.
- Lakretz, Yair, Théo Desbordes, Dieuwke Hupkes, and Stanislas Dehaene. 2022. Can transformers process recursive nested constructions, like humans? In *Proceedings of the 29th International Conference on Computational Linguistics*, pages 3226–3232, International Committee on Computational Linguistics, Gyeongju, Republic of Korea.
- Lal, Yash Kumar, Niket Tandon, Tanvi Aggarwal, Horace Liu, Nathanael Chambers, Raymond Mooney, and Niranjan Balasubramanian. 2022. Using commonsense knowledge to answer why-questions. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 1204–1219.
- Lampinen, Andrew. 2022. Can language models handle recursively nested grammatical structures? a case study on comparing models and humans. *ArXiv*, arXiv:2210.15303.
- Lampinen, Andrew, Ishita Dasgupta, Stephanie Chan, Kory Mathewson, Mh Tessler, Antonia Creswell, James McClelland, Jane Wang, and Felix Hill. 2022. Can language models learn from explanations in context? In *Findings of the Association for Computational Linguistics: EMNLP 2022*, pages 537–563.
- Lasri, Karim, Alessandro Lenci, and Thierry Poibeau. 2022a. Does BERT really agree ? fine-grained analysis of lexical dependence on a syntactic task. In *Findings of the Association for Computational Linguistics: ACL 2022*, pages 2309–2315.
- Lasri, Karim, Alessandro Lenci, and Thierry Poibeau. 2022b. Word order matters when you increase masking. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 1808–1815.
- Lasri, Karim, Olga Seminck, Alessandro Lenci, and Thierry Poibeau. 2022. Subject verb agreement error patterns in meaningless sentences: Humans vs. BERT. In *Proceedings of the 29th International Conference on Computational Linguistics*, pages 37–43, International Committee on Computational Linguistics, Gyeongju,

- Republic of Korea.
- Lee, Angie. 2023. What are large language models used for? *NVIDIA Blog*.
- Lee, Jooyoung, Thai Le, Jinghui Chen, and Dongwon Lee. 2023. Do language models plagiarize? In *The ACM Web Conference*.
- Lee, Katherine, Daphne Ippolito, Andrew Nystrom, Chiyuan Zhang, Douglas Eck, Chris Callison-Burch, and Nicholas Carlini. 2022. Deduplicating training data makes language models better. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 8424–8445.
- Lee, Nayeon, Yejin Bang, Andrea Madotto, and Pascale Fung. 2021. Towards few-shot fact-checking via perplexity. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 1971–1981.
- Lee, Soo-Hwan and Sebastian Schuster. 2022. Can language models capture syntactic associations without surface cues? a case study of reflexive anaphor licensing in English control constructions. In *Proceedings of the Society for Computation in Linguistics 2022*, pages 206–211.
- Lehman, Eric, Sarthak Jain, Karl Pichotta, Yoav Goldberg, and Byron Wallace. 2021. Does BERT pretrained on clinical notes reveal sensitive data? In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 946–959.
- Lester, Brian, Rami Al-Rfou, and Noah Constant. 2021. The power of scale for parameter-efficient prompt tuning. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 3045–3059.
- Levy, Sharon, Emily Allaway, Melanie Subbiah, Lydia Chilton, Desmond Patton, Kathleen McKeown, and William Yang Wang. 2022. SafeText: A benchmark for exploring physical safety in language models. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 2407–2421.
- Levy, Sharon, Michael Saxon, and William Yang Wang. 2021. Investigating memorization of conspiracy theories in text generation. In *Findings of the Association for Computational Linguistics: ACL-IJCNLP 2021*, pages 4718–4729.
- Li, Jiaxuan, Lang Yu, and Allyson Ettinger. 2022. Counterfactual reasoning: Do language models need world knowledge for causal inference? In *Workshop on Neuro Causal and Symbolic AI (nCSI)*.
- Li, Xiang Lorraine, Adhiguna Kuncoro, Jordan Hoffmann, Cyprien de Masson d’Autume, Phil Blunsom, and Aida Nematzadeh. 2022a. A systematic investigation of commonsense knowledge in large language models. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 11838–11855.
- Li, Xingxuan, Yutong Li, Linlin Liu, Lidong Bing, and Shafiq Joty. 2022b. Is GPT-3 a psychopath? evaluating large language models from a psychological perspective. *ArXiv*, arXiv:2212.10529.
- Lieber, Opher, Or Sharir, Barak Lenz, and Yoav Shoham. 2021. Jurassic-1: Technical details and evaluation. *White Paper. AI21 Labs*.
- Lin, Bill Yuchen, Seyeon Lee, Rahul Khanna, and Xiang Ren. 2020. Birds have four legs?! NumerSense: Probing Numerical Commonsense Knowledge of Pre-Trained Language Models. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 6862–6868.
- Lin, Stephanie, Jacob Hilton, and Owain Evans. 2022. TruthfulQA: Measuring how models mimic human falsehoods. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 3214–3252.
- Liu, Emmy, Chenxuan Cui, Kenneth Zheng, and Graham Neubig. 2022a. Testing the ability of language models to interpret figurative language. In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 4437–4452.
- Liu, Fangyu, Julian Eisenschlos, Jeremy Cole, and Nigel Collier. 2022b. Do ever larger octopi still amplify reporting biases? evidence from judgments of typical colour. In *Proceedings of the 2nd Conference of the Asia-Pacific Chapter of the Association for Computational Linguistics and the 12th International Joint Conference on Natural Language Processing (Volume 2: Short Papers)*, pages 210–220.
- Liu, Ruibo, Chenyan Jia, Jason Wei, Guangxuan Xu, and Soroush Vosoughi. 2022c. Quantifying and alleviating political bias in language models. *Artificial Intelligence*, 304:103654.
- Liu, Yinhan, Myle Ott, Naman Goyal, Jingfei Du, Mandar Joshi, Danqi Chen, Omer

- Levy, Mike Lewis, Luke Zettlemoyer, and Veselin Stoyanov. 2019. RoBERTa: A robustly optimized BERT pretraining approach. *arXiv*, arXiv:1907.11692.
- Liu, Zeyu, Yizhong Wang, Jungo Kasai, Hannaneh Hajishirzi, and Noah A. Smith. 2021. Probing across time: What does RoBERTa know and when? In *Findings of the Association for Computational Linguistics: EMNLP 2021*, pages 820–842.
- Magee, Liam, Lida Ghahremanlou, Karen Soldatic, and Shanthi Robertson. 2021. Intersectional bias in causal language models. *ArXiv*, arXiv:2107.07691.
- Mahowald, Kyle. 2023. A discerning several thousand judgments: GPT-3 rates the article + adjective + numeral + noun construction. In *Proceedings of the Conference of the European Chapter of the Association for Computational Linguistics*.
- Mahowald, Kyle, Evgeniia Diachek, Edward Gibson, Evelina Fedorenko, and Richard Futrell. 2022. Experimentally measuring the redundancy of grammatical cues in transitive clauses. *ArXiv*, arXiv:2201.12911.
- Malkin, Nikolay, Sameera Lanka, Pranav Goel, and Nebojsa Jojic. 2021. Studying word order through iterative shuffling. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 10351–10366.
- Mallen, Alex, Akari Asai, Victor Zhong, Rajarshi Das, Hannaneh Hajishirzi, and Daniel Khashabi. 2022. When not to trust language models: Investigating effectiveness and limitations of parametric and non-parametric memories. *ArXiv*, arXiv:2212.10511.
- Marr, David. 2010. *Vision: A Computational Investigation into the Human Representation and Processing of Visual Information*. The MIT Press.
- Masis, Tessa and Carolyn Anderson. 2021. ProSPer: Probing human and neural network language model understanding of spatial perspective. In *Proceedings of the Fourth BlackboxNLP Workshop on Analyzing and Interpreting Neural Networks for NLP*, pages 95–135.
- Massarelli, Luca, Fabio Petroni, Aleksandra Piktus, Myle Ott, Tim Rocktaschel, Vassilis Plachouras, Fabrizio Silvestri, and Sebastian Riedel. 2020. How decoding strategies affect the verifiability of generated text. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, pages 223–235.
- Mattern, Justus, Zhijing Jin, Mrinmaya Sachan, Rada Mihalcea, and B. Schölkopf. 2022. Understanding stereotypes in language models: Towards robust measurement and zero-shot debiasing. *ArXiv*, arXiv:2212.10678.
- McCoy, R. Thomas, Paul Smolensky, Tal Linzen, Jianfeng Gao, and Aslı Celikyilmaz. 2021. How much do language models copy from their training data? evaluating linguistic novelty in text generation using RAVEN. *ArXiv*, arXiv:2111.09509.
- McCoy, Thomas, Robert Frank, and Tal Linzen. 2018. Revisiting the poverty of the stimulus: Hierarchical generalization without a hierarchical bias in recurrent neural networks. In *Proceedings of the Annual Meeting of the Cognitive Science Society*, volume 40.
- McCoy, Tom, Ellie Pavlick, and Tal Linzen. 2019. Right for the wrong reasons: Diagnosing syntactic heuristics in natural language inference. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 3428–3448.
- Meade, Nicholas, Elinor Poole-Dayana, and Siva Reddy. 2022. An empirical survey of the effectiveness of debiasing techniques for pre-trained language models. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1878–1898.
- Mehrabi, Ninareh, Ahmad Beirami, Fred Morstatter, and Aram Galstyan. 2022. Robust conversational agents against imperceptible toxicity triggers. In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 2831–2847.
- Meister, Clara and Ryan Cotterell. 2021. Language model evaluation beyond perplexity. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 5328–5339.
- Meng, Kevin, David Bau, Alex J Andonian, and Yonatan Belinkov. 2022. Locating and editing factual associations in GPT. In *Advances in Neural Information Processing Systems*.
- Miaschi, Alessio, Dominique Brunato, Felice Dell’Orletta, and Giulia Venturi. 2020. Linguistic profiling of a neural language model. In *Proceedings of the 28th International Conference on Computational Linguistics*, pages 745–756, International

- Committee on Computational Linguistics, Barcelona, Spain (Online).
- Michaelov, James and Benjamin Bergen. 2022a. Collateral facilitation in humans and language models. In *Proceedings of the 26th Conference on Computational Natural Language Learning (CoNLL)*, pages 13–26.
- Michaelov, James and Benjamin Bergen. 2022b. ‘rarely’ a problem? language models exhibit inverse scaling in their predictions following ‘few’-type quantifiers. *ArXiv*, arXiv:2212.08700.
- Min, Sewon, Xixi Lyu, Ari Holtzman, Mikel Artetxe, Mike Lewis, Hannaneh Hajishirzi, and Luke Zettlemoyer. 2022. Rethinking the role of demonstrations: What makes in-context learning work? In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 11048–11064.
- Miotto, Marilù, Nicola Rossberg, and Bennett Kleinberg. 2022. Who is GPT-3? an exploration of personality, values and demographics. In *Proceedings of the Fifth Workshop on Natural Language Processing and Computational Social Science (NLP+CSS)*, pages 218–227.
- Misra, Kanishka. 2022. minicons: Enabling flexible behavioral and representational analyses of Transformer language models. *ArXiv*, arXiv:2203.13112.
- Misra, Kanishka, Allyson Ettinger, and Julia Rayz. 2020. Exploring BERT’s sensitivity to lexical cues using tests from semantic priming. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, pages 4625–4635.
- Misra, Kanishka, Allyson Ettinger, and Julia Taylor Rayz. 2021. Do language models learn typicality judgments from text? In *Proceedings of the Annual Meeting of the Cognitive Science Society*, volume 43.
- Misra, Kanishka, Julia Taylor Rayz, and Allyson Ettinger. 2023. COMPS: Conceptual minimal pair sentences for testing property knowledge and inheritance in pre-trained language models. In *Proceedings of the Conference of the European Chapter of the Association for Computational Linguistics*.
- Mitchell, Melanie and David Krakauer. 2022. The debate over understanding in AI’s large language models. *ArXiv*, arXiv:2210.13966.
- Monroe, Will and Christopher Potts. 2015. Learning in the rational speech acts model. *arXiv*, arXiv:1510.06807.
- Mosbach, Marius, Anna Khokhlova, Michael A. Hedderich, and Dietrich Klakow. 2020. On the interplay between fine-tuning and sentence-level probing for linguistic knowledge in pre-trained Transformers. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, pages 2502–2516.
- Nadeem, Moin, Anna Bethke, and Siva Reddy. 2021. StereoSet: Measuring stereotypical bias in pretrained language models. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 5356–5371.
- Nangia, Nikita, Clara Vania, Rasika Bhalerao, and Samuel R. Bowman. 2020. CrowS-pairs: A challenge dataset for measuring social biases in masked language models. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 1953–1967.
- Nayak, Pandu. 2019. Understanding searches better than ever before. *The Keyword: Google Blog*.
- Newman, Benjamin, Kai-Siang Ang, Julia Gong, and John Hewitt. 2021. Refining targeted syntactic evaluation of language models. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 3710–3723.
- Nozza, Debora, Federico Bianchi, and Dirk Hovy. 2021. HONEST: Measuring hurtful sentence completion in language models. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 2398–2406.
- Nozza, Debora, Federico Bianchi, Anne Lauscher, and Dirk Hovy. 2022. Measuring harmful sentence completion in language models for LGBTQIA+ individuals. In *Proceedings of the Second Workshop on Language Technology for Equality, Diversity and Inclusion*, pages 26–34.
- O’Connor, Joe and Jacob Andreas. 2021. What context features can transformer language models use? In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 851–864.
- Olsson, Catherine, Nelson Elhage, Neel Nanda, Nicholas Joseph, Nova DasSarma, Tom Henighan, Benjamin Mann, Amanda Askell, Yuntao Bai, Anna Chen, Tom

- Conerly, Dawn Drain, Deep Ganguli, Zac Hatfield-Dodds, Danny Hernandez, Scott Johnston, Andy Jones, John Kernion, Liane Lovitt, Kamal Ndousse, Dario Amodei, Tom B. Brown, Jack Clark, Jared Kaplan, Sam McCandlish, and Chris Olah. 2022. In-context learning and induction heads. *ArXiv*, arXiv:2209.11895.
- OpenAI. 2022. ChatGPT: Optimizing language models for dialogue. *OpenAI Blog*.
- OpenAI. 2023a. GPT-4 technical report. *OpenAI*.
- OpenAI. 2023b. Model index for researchers. *OpenAI*.
- Ousidhoum, Nedjma, Xinran Zhao, Tianqing Fang, Yangqiu Song, and Dit-Yan Yeung. 2021. Probing toxic content in large pre-trained language models. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 4262–4274.
- Ouyang, Long, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Gray, John Schulman, Jacob Hilton, Fraser Kelton, Luke Miller, Maddie Simens, Amanda Askell, Peter Welinder, Paul Christiano, Jan Leike, and Ryan Lowe. 2022. Training language models to follow instructions with human feedback. In *Advances in Neural Information Processing Systems*.
- Paik, Cory, Stéphane Aroca-Ouellette, Alessandro Roncone, and Katharina Kann. 2021. The World of an Octopus: How Reporting Bias Influences a Language Model’s Perception of Color. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 823–835.
- Pandia, Lalchand, Yan Cong, and Allyson Ettinger. 2021. Pragmatic competence of pre-trained language models through the lens of discourse connectives. In *Proceedings of the 25th Conference on Computational Natural Language Learning*, pages 367–379.
- Pandia, Lalchand and Allyson Ettinger. 2021. Sorting through the noise: Testing robustness of information processing in pre-trained language models. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 1583–1596.
- Pandit, Onkar and Yufang Hou. 2021. Probing for bridging inference in transformer language models. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 4153–4163.
- Park, Kwonsik, Myung-Kwan Park, and Sanghoun Song. 2021. Deep learning can contrast the minimal pairs of syntactic data. *Linguistic Research*, 38(2):395–424.
- Patel, Roma and Ellie Pavlick. 2021. “was it “stated” or was it “claimed”? : How linguistic bias affects generative language models. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 10080–10095.
- Pedinotti, Paolo, Eliana Di Palma, Ludovica Cerini, and Alessandro Lenci. 2021a. A howling success or a working sea? testing what BERT knows about metaphors. In *Proceedings of the Fourth BlackboxNLP Workshop on Analyzing and Interpreting Neural Networks for NLP*, pages 192–204.
- Pedinotti, Paolo, Giulia Rambelli, Emmanuele Chersoni, Enrico Santus, Alessandro Lenci, and Philippe Blache. 2021b. Did the cat drink the coffee? challenging transformers with generalized event knowledge. In *Proceedings of *SEM 2021: The Tenth Joint Conference on Lexical and Computational Semantics*, pages 1–11.
- Peng, Hao, Xiaozhi Wang, Shengding Hu, Hailong Jin, Lei Hou, Juanzi Li, Zhiyuan Liu, and Qun Liu. 2022. COPEN: Probing conceptual knowledge in pre-trained language models. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 5015–5035.
- Penha, Gustavo and Claudia Hauff. 2020. What does BERT know about books, movies and music? probing BERT for conversational recommendation. In *Proceedings of the 14th ACM Conference on Recommender Systems*, page 388–397, Association for Computing Machinery, New York, NY, USA.
- Perez, Ethan, Saffron Huang, Francis Song, Trevor Cai, Roman Ring, John Aslanides, Amelia Glaese, Nat McAleese, and Geoffrey Irving. 2022a. Red teaming language models with language models. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 3419–3448.
- Perez, Ethan, Sam Ringer, Kamile Lukosiute, Karina Nguyen, Edwin Chen, Scott Heiner, Craig Pettit, Catherine Olsson, Sandipan Kundu, Saurav Kadavath, Andy Jones, Anna Chen, Benjamin Mann, Brian Israel, Bryan Seethor, Cameron McKinnon,

- Christopher Olah, Daisong Yan, Daniela Amodei, Dario Amodei, Dawn Drain, Dustin Li, Eli Tran-Johnson, Guro Khundadze, John Kernion, James Landis, Jamie Kerr, Jared Mueller, Jeeyoon Hyun, Joshua Landau, Kamal Ndousse, Landon Goldberg, Liane Lovitt, Martin Lucas, Michael Sellitto, Miranda Zhang, Neerav Kingsland, Nelson Elhage, Nicholas Joseph, Noemí Mercado, Nova DasSarma, Oliver Rausch, Robin Larson, Sam McCandlish, Scott Johnston, Shauna Kravec, Sheer El Showk, Tamera Lanham, Timothy Telleen-Lawton, Tom Brown, Tom Henighan, Tristan Hume, Yuntao Bai, Zac Hatfield-Dodds, Jack Clark, Sam Bowman, Amanda Askell, Roger Grosse, Danny Hernandez, Deep Ganguli, Evan Hubinger, Nicholas Schiefer, and Jared Kaplan. 2022b. Discovering language model behaviors with model-written evaluations. *ArXiv*, arXiv:2212.09251.
- Pérez-Mayos, Laura, Miguel Ballesteros, and Leo Wanner. 2021. How much pretraining data do language models need to learn syntax? In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 1571–1582.
- Petroni, Fabio, Tim Rocktäschel, Sebastian Riedel, Patrick Lewis, Anton Bakhtin, Yuxiang Wu, and Alexander Miller. 2019. Language models as knowledge bases? In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 2463–2473.
- Petty, Jackson and Robert Frank. 2021. Transformers generalize linearly. *ArXiv*, arXiv:2109.12036.
- Piantadosi, Steven and Evelina Fedorenko. 2017. Infinitely productive language can arise from chance under communicative pressure. *Journal of Language Evolution*, 2(2):141–147.
- Podkorytov, Maksim, Daniel Bis, and Xiuwen Liu. 2021. How can the [MASK] know? the sources and limitations of knowledge in BERT. In *IEEE International Joint Conference on Neural Networks*.
- Poerner, Nina, Ulli Waltinger, and Hinrich Schütze. 2019. BERT is not a knowledge base (yet): Factual knowledge vs. name-based reasoning in unsupervised QA. *ArXiv*, arXiv:1911.03681.
- Porada, Ian, Alessandro Sordani, and Jackie Cheung. 2022. Does pre-training induce systematic inference? how masked language models acquire commonsense knowledge. In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 4550–4557.
- Prefors, Amy, Terry Regier, and Joshua Tenenbaum. 2006. Poverty of the stimulus? a rational approach. In *Proceedings of the Annual Meeting of the Cognitive Science Society*, volume 28.
- Press, Ofir, Noah Smith, and Mike Lewis. 2022. Train short, test long: Attention with linear biases enables input length extrapolation. In *International Conference on Learning Representations*.
- Press, Ofir, Muru Zhang, Sewon Min, Ludwig Schmidt, Noah A. Smith, and Mike Lewis. 2022. Measuring and narrowing the compositionality gap in language models. *ArXiv*, arXiv:2210.03350.
- Qin, Lianhui, Aditya Gupta, Shyam Upadhyay, Luheng He, Yejin Choi, and Manaal Faruqui. 2021. TIMEDIAL: Temporal commonsense reasoning in dialog. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 7066–7076.
- Qiu, Linlu, Peter Shaw, Panupong Pasupat, Tianze Shi, Jonathan Herzig, Emily Pitler, Fei Sha, and Kristina Toutanova. 2022. Evaluating the impact of model scale for compositional generalization in semantic parsing. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 9157–9179.
- Radford, Alec, Jong Wook Kim, Tao Xu, Greg Brockman, Christine McLeavey, and Ilya Sutskever. 2022. Robust speech recognition via large-scale weak supervision. *arXiv*, arXiv:2212.04356.
- Radford, Alec, Karthik Narasimhan, Tim Salimans, and Ilya Sutskever. 2018. Improving language understanding by generative pre-training. *OpenAI*.
- Radford, Alec, Jeff Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. 2019. Language models are unsupervised multitask learners. *OpenAI*.
- Rae, Jack W., Sebastian Borgeaud, Trevor Cai, Katie Millican, Jordan Hoffmann, Francis Song, John Aslanides, Sarah Henderson, Roman Ring, Susannah Young, Eliza Rutherford, Tom Hennigan, Jacob Menick, Albin Cassirer, Richard Powell, George van den Driessche, Lisa Anne Hendricks, Maribeth Rauh, Po-Sen Huang, Amelia Glaese, Johannes Welbl, Sumanth

- Dathathri, Saffron Huang, Jonathan Uesato, John F. J. Mellor, Irina Higgins, Antonia Creswell, Nathan McAleese, Amy Wu, Erich Elsen, Siddhant M. Jayakumar, Elena Buchatskaya, David Budden, Esme Sutherland, Karen Simonyan, Michela Paganini, L. Sifre, Lena Martens, Xiang Lorraine Li, Adhiguna Kuncoro, Aida Nematzadeh, Elena Gribovskaya, Domenic Donato, Angeliki Lazaridou, Arthur Mensch, Jean-Baptiste Lespiau, Maria Tsimpoukelli, N. K. Grigorev, Doug Fritz, Thibault Sottiaux, Mantas Pajarskas, Tobias Pohlen, Zhitao Gong, Daniel Toyama, Cyprien de Masson d’Autume, Yujia Li, Tayfun Terzi, Vladimir Mikulik, Igor Babuschkin, Aidan Clark, Diego de Las Casas, Aurelia Guy, Chris Jones, James Bradbury, Matthew G. Johnson, Blake A. Hechtman, Laura Weidinger, Iason Gabriel, William S. Isaac, Edward Lockhart, Simon Osindero, Laura Rimell, Chris Dyer, Oriol Vinyals, Kareem W. Ayoub, Jeff Stanway, L. L. Bennett, Demis Hassabis, Koray Kavukcuoglu, and Geoffrey Irving. 2021. Scaling language models: Methods, analysis & insights from training gopher. *ArXiv*, arXiv:2112.11446.
- Raffel, Colin, Noam Shazeer, Adam Roberts, Katherine Lee, Sharan Narang, Michael Matena, Yanqi Zhou, Wei Li, and Peter J. Liu. 2020. Exploring the limits of transfer learning with a unified text-to-text Transformer. *Journal of Machine Learning Research*, 21(140).
- Raj, Harsha, Domenic Rosati, and Subhabrata Majumdar. 2022. Measuring reliability of large language models through semantic consistency. In *NeurIPS ML Safety Workshop*.
- Ravichander, Abhilasha, Eduard Hovy, Kaheer Suleman, Adam Trischler, and Jackie Chi Kit Cheung. 2020. On the systematicity of probing contextualized word representations: The case of hypernymy in BERT. In *Proceedings of the Ninth Joint Conference on Lexical and Computational Semantics*, pages 88–102.
- Razeghi, Yasaman, Robert Logan IV, Matt Gardner, and Sameer Singh. 2022. Impact of pretraining term frequencies on few-shot numerical reasoning. In *Findings of the Association for Computational Linguistics: EMNLP 2022*, pages 840–854.
- Reif, Emily, Daphne Ippolito, Ann Yuan, Andy Coenen, Chris Callison-Burch, and Jason Wei. 2022. A recipe for arbitrary text style transfer with large language models. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, pages 837–848.
- Reif, Emily, Ann Yuan, Martin Wattenberg, Fernanda B Viegas, Andy Coenen, Adam Pearce, and Been Kim. 2019. Visualizing and measuring the geometry of BERT. In *Advances in Neural Information Processing Systems*, volume 32.
- Rogers, Anna, Olga Kovaleva, and Anna Rumshisky. 2020. A primer in BERTology: What we know about how BERT works. *Transactions of the Association for Computational Linguistics*, 8:842–866.
- Romero, Julien and Simon Razniewski. 2022. Do children texts hold the key to commonsense knowledge? In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 10954–10959.
- Ruis, Laura, Akbir Khan, Stella Rose Biderman, Sara Hooker, Tim Rocktaschel, and Edward Grefenstette. 2022. Large language models are not zero-shot communicators. *ArXiv*, arXiv:2210.14986.
- Ryu, Soo Hyun and Richard Lewis. 2021. Accounting for agreement phenomena in sentence comprehension with transformer language models: Effects of similarity-based interference on surprisal and attention. In *Proceedings of the Workshop on Cognitive Modeling and Computational Linguistics*, pages 61–71.
- Sahu, Pritish, Michael Cogswell, Yunye Gong, and Ajay Divakaran. 2022. Unpacking large language models with conceptual consistency. *ArXiv*, arXiv:2209.15093.
- Sancheti, Abhilasha and Rachel Rudinger. 2022. What do large language models learn about scripts? In *Proceedings of the 11th Joint Conference on Lexical and Computational Semantics*, pages 1–11.
- Sanh, Victor, Lysandre Debut, Julien Chaumond, and Thomas Wolf. 2019. DistilBERT, a distilled version of BERT: Smaller, faster, cheaper and lighter. In *Workshop on Energy Efficient Machine Learning and Cognitive Computing*.
- Sap, Maarten, Ronan Le Bras, Daniel Fried, and Yejin Choi. 2022. Neural theory-of-mind? on the limits of social intelligence in large LMs. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 3762–3780.
- Saparov, Abulhair and He He. 2023. Language models are greedy reasoners: A systematic formal analysis of chain-of-thought. In *International*

- Conference on Learning Representations. van Schijndel, Marten, Aaron Mueller, and Tal Linzen. 2019. Quantity doesn't buy quality syntax with neural language models. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 5831–5837.
- Schuster, Sebastian and Tal Linzen. 2022. When a sentence does not introduce a discourse entity, transformer-based models still sometimes refer to it. In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 969–982.
- Senel, Lutfi Kerem and Hinrich Schütze. 2021. Does she wink or does she nod? a challenging benchmark for evaluating word understanding of language models. In *Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics: Main Volume*, pages 532–538.
- Sennrich, Rico, Barry Haddow, and Alexandra Birch. 2016. Neural machine translation of rare words with subword units. In *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1715–1725.
- Serrano, Sofia and Noah A. Smith. 2019. Is attention interpretable? In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 2931–2951.
- Seshadri, Preethi, Pouya Pezeshkpour, and Sameer Singh. 2022. Quantifying social biases using templates is unreliable. In *Workshop on Trustworthy and Socially Responsible Machine Learning*.
- Shaikh, Omar, Hongxin Zhang, William B. Held, Michael Bernstein, and Diyi Yang. 2022. On second thought, let's not think step by step! Bias and toxicity in zero-shot reasoning. *ArXiv*, arXiv:2212.08061.
- Shardlow, Matthew and Piotr Przybyla. 2022. Deanthropomorphising NLP: Can a language model be conscious? *ArXiv*, arXiv:2211.11483.
- Shaw, Peter, Jakob Uszkoreit, and Ashish Vaswani. 2018. Self-attention with relative position representations. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 2 (Short Papers)*, pages 464–468.
- Sheng, Emily, Kai-Wei Chang, Prem Natarajan, and Nanyun Peng. 2021a. "nice try, kiddo": Investigating ad hominem in dialogue responses. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 750–767.
- Sheng, Emily, Kai-Wei Chang, Prem Natarajan, and Nanyun Peng. 2021b. Societal biases in language generation: Progress and challenges. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 4275–4293.
- Sheng, Emily, Kai-Wei Chang, Premkumar Natarajan, and Nanyun Peng. 2019. The woman worked as a babysitter: On biases in language generation. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 3407–3412.
- Shi, Freda, Xinyun Chen, Kanishka Misra, Nathan Scales, David Dohan, Ed Chi, Nathanael Schärli, and Denny Zhou. 2023. Large language models can be easily distracted by irrelevant context. *arXiv*, arXiv:2302.00093.
- Shi, Haohan and Phillip Wolff. 2021. What Transformers might know about the physical world: T5 and the origins of knowledge. In *Proceedings of the Annual Meeting of the Cognitive Science Society*, volume 43.
- Shwartz, Vered and Yejin Choi. 2020. Do neural language models overcome reporting bias? In *Proceedings of the 28th International Conference on Computational Linguistics*, pages 6863–6870, International Committee on Computational Linguistics, Barcelona, Spain (Online).
- Shwartz, Vered, Rachel Rudinger, and Øyvind Tafjord. 2020. "you are grounded!": Latent name artifacts in pre-trained language models. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 6850–6861.
- Sikos, Les, Noortje Venhuizen, Heiner Drenhaus, and Matthew Crocker. 2021. Reevaluating pragmatic reasoning in language games. *PLOS One*, 16(3).
- Silva, Andrew, Pradyumna Tambwekar, and Matthew Gombolay. 2021. Towards a comprehensive understanding and

- accurate evaluation of societal biases in pre-trained transformers. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 2383–2389.
- Simmons, Gabriel. 2022. Moral mimicry: Large language models produce moral rationalizations tailored to political identity. *ArXiv*, arXiv:2209.12106.
- Sinclair, Arabella, Jaap Jumelet, Willem Zuidema, and Raquel Fernández. 2022. Structural persistence in language models: Priming as a window into abstract language representations. *Transactions of the Association for Computational Linguistics*, 10:1031–1050.
- Sinha, Koustuv, Jon Gauthier, Aaron Mueller, Kanishka Misra, Keren Fuentes, Roger Levy, and Adina Williams. 2022a. Language model acceptability judgements are not always robust to context. *ArXiv*, arXiv:2212.08979.
- Sinha, Koustuv, Robin Jia, Dieuwke Hupkes, Joelle Pineau, Adina Williams, and Douwe Kiela. 2021. Masked language modeling and the distributional hypothesis: Order word matters pre-training for little. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 2888–2913.
- Sinha, Koustuv, Amirhossein Kazemnejad, Siva Reddy, Joelle Pineau, Dieuwke Hupkes, and Adina Williams. 2022b. The curious case of absolute position embeddings. In *Findings of the Association for Computational Linguistics: EMNLP 2022*, pages 4449–4472.
- Smith, Eric Michael, Melissa Hall, Melanie Kambadur, Eleonora Presani, and Adina Williams. 2022a. “I’m sorry to hear that”: Finding new biases in language models with a holistic descriptor dataset. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 9180–9211.
- Smith, Shaden, Mostofa Patwary, Brandon Norick, Patrick LeGresley, Samyam Rajbhandari, Jared Casper, Zhun Liu, Shrimai Prabhumoye, George Zerveas, Vijay Anand Korthikanti, Elton Zhang, Rewon Child, Reza Yazdani Aminabadi, Julie Bernauer, Xia Song, Mohammad Shoeybi, Yuxiong He, Michael Houston, Saurabh Tiwary, and Bryan Catanzaro. 2022b. Using DeepSpeed and Megatron to train Megatron-Turing NLG 530B, a large-scale generative language model. *arXiv*, arXiv:2201.11990.
- Spitale, Giovanni, Nikola Biller-Andorno, and Federico Germani. 2023. AI model GPT-3 (dis)informs us better than humans. *ArXiv*, arXiv:2301.11924.
- Srivastava, Aarohi, Abhinav Rastogi, Abhishek Rao, Abu Awal Md Shoeb, Abubakar Abid, Adam Fisch, Adam R Brown, Adam Santoro, Aditya Gupta, Adrià Garriga-Alonso, et al. 2022. Beyond the imitation game: Quantifying and extrapolating the capabilities of language models. *arXiv*, arXiv:2206.04615.
- Stammbach, Dominik, Maria Antoniak, and Elliott Ash. 2022. Heroes, villains, and victims, and GPT-3: Automated extraction of character roles without training data. In *Proceedings of the 4th Workshop of Narrative Understanding (WNU2022)*, pages 47–56.
- Stevenson, Claire, Iris Smal, Matthijs Baas, Raoul Grasman, and Han van der Maas. 2022. Putting GPT-3’s creativity to the (alternative uses) test. In *International Conference on Computational Creativity*.
- Stolfo, Alessandro, Zhijing Jin, Kumar Shridhar, Bernhard Schölkopf, and Mrinmaya Sachan. 2022. A causal framework to quantify the robustness of mathematical reasoning with language models. *ArXiv*, arXiv:2210.12023.
- Strubell, Emma, Ananya Ganes, and Andrew McCallum. 2019. Energy and policy considerations for deep learning in NLP. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 3645–3650.
- Su, Jianlin, Yu Lu, Shengfeng Pan, Ahmed Murtadha, Bo Wen, and Yunfeng Liu. 2021. RoFormer: Enhanced Transformer with rotary position embedding. *arXiv*, arXiv:2104.09864.
- Summers-Stay, Douglas, Claire Bonial, and Clare Voss. 2021. What can a generative language model answer about a passage? In *Proceedings of the 3rd Workshop on Machine Reading for Question Answering*, pages 73–81.
- Suzgun, Mirac, Nathan Scales, Nathanael Scharli, Sebastian Gehrmann, Yi Tay, Hyung Won Chung, Aakanksha Chowdhery, Quoc V. Le, Ed Chi, Denny Zhou, and Jason Wei. 2022. Challenging BIG-Bench tasks and whether chain-of-thought can solve them. *ArXiv*, arXiv:2210.09261.
- Swamy, Vinitra, Angelika Romanou, and Martin Jaggi. 2021. Interpreting language models through knowledge graph extraction. In *Workshop on eXplainable AI Approaches for Debugging and Diagnosis*.

- Tal, Yarden, Inbal Magar, and Roy Schwartz. 2022. Fewer errors, but more stereotypes? the effect of model size on gender bias. In *Proceedings of the 4th Workshop on Gender Bias in Natural Language Processing (GeBNLP)*, pages 112–120.
- Tamborrino, Alexandre, Nicola Pellicanò, Baptiste Pannier, Pascal Voitot, and Louise Naudin. 2020. Pre-training is (almost) all you need: An application to commonsense reasoning. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 3878–3887.
- Tang, Kenan and Hanchun Jiang. 2022. Gender biases unexpectedly fluctuate in the pre-training stage of masked language models. *ArXiv*, arXiv:2211.14639.
- Tay, Yi, Mostafa Dehghani, Samira Abnar, Hyung Won Chung, W. Fedus, Jinfeng Rao, Sharan Narang, Vinh Tran, Dani Yogatama, and Donald Metzler. 2022a. Scaling laws vs model architectures: How does inductive bias influence scaling? *ArXiv*, arXiv:2207.10551.
- Tay, Yi, Mostafa Dehghani, Jinfeng Rao, William Fedus, Samira Abnar, Hyung Won Chung, Sharan Narang, Dani Yogatama, Ashish Vaswani, and Donald Metzler. 2022b. Scale efficiently: Insights from pre-training and fine-tuning Transformers. In *International Conference on Learning Representations*.
- Tejada, Giorgia Nidia Carranza, Johannes Scholtes, and Gerasimos Spanakis. 2021. A study of BERT’s processing of negations to determine sentiment. In *Benelux Conference on Artificial Intelligence and the Belgian Dutch Conference on Machine Learning*.
- Tenney, Ian, Dipanjan Das, and Ellie Pavlick. 2019. BERT rediscovers the classical NLP pipeline. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 4593–4601.
- Thewsey, Alex. 2021. Bring structure to diverse documents with Amazon Textract and transformer-based models on Amazon SageMaker. *AWS Machine Learning Blog*.
- Tirumala, Kushal, Aram H. Markosyan, Luke Zettlemoyer, and Armen Aghajanyan. 2022. Memorization without overfitting: Analyzing the training dynamics of large language models. In *Advances in Neural Information Processing Systems*.
- Touileb, Samia. 2022. Exploring the effects of negation and grammatical tense on bias probes. In *Proceedings of the 2nd Conference of the Asia-Pacific Chapter of the Association for Computational Linguistics and the 12th International Joint Conference on Natural Language Processing (Volume 2: Short Papers)*, pages 423–429.
- Traylor, Aaron, Roman Feiman, and Ellie Pavlick. 2021. AND does not mean OR: Using formal languages to study language models’ representations. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 2: Short Papers)*, pages 158–167.
- Trott, Sean. 2023. In cautious defense of LLM-ology. *Blog Post*.
- Trott, Sean, Cameron J. Jones, Tyler A. Chang, James Michaelov, and Benjamin Bergen. 2022. Do large language models know what humans know? *ArXiv*, arXiv:2209.01515.
- Truong, Thinh Hung, Yulia Otmakhova, Timothy Baldwin, Trevor Cohn, Jey Han Lau, and Karin Verspoor. 2022. Not another negation benchmark: The NaN-NLI test suite for sub-clausal negation. In *Proceedings of the 2nd Conference of the Asia-Pacific Chapter of the Association for Computational Linguistics and the 12th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 883–894.
- Tuckute, Greta, Aalok Sathe, Mingye Wang, Harley Yoder, Cory Shain, and Evelina Fedorenko. 2022. SentSpace: Large-scale benchmarking and evaluation of text using cognitively motivated lexical, syntactic, and semantic features. In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies: System Demonstrations*, pages 99–113.
- Upadhye, Shiva, Leon Bergen, and Andrew Kehler. 2020. Predicting reference: What do language models learn about discourse models? In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 977–982.
- Ushio, Asahi, Luis Espinosa Anke, Steven Schockaert, and Jose Camacho-Collados. 2021. BERT is to NLP what AlexNet is to CV: Can pre-trained language models identify analogies? In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 3609–3624.

- Valmeekam, Karthik, Alberto Olmo, Sarath Sreedharan, and Subbarao Kambhampati. 2022. Large language models still can't plan (a benchmark for LLMs on planning and reasoning about change). In *Foundation Models for Decision Making Workshop*.
- de Vassimon Manela, Daniel, David Errington, Thomas Fisher, Boris van Breugel, and Pasquale Minervini. 2021. Stereotype and skew: Quantifying gender bias in pre-trained and fine-tuned language models. In *Proceedings of the 16th Conference of the European Chapter of the Association for Computational Linguistics: Main Volume*, pages 2232–2242.
- Vaswani, Ashish, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan Gomez, Łukasz Kaiser, and Illia Polosukhin. 2017. Attention is all you need. In *Advances in Neural Information Processing Systems*, volume 30.
- Vig, Jesse and Yonatan Belinkov. 2019. Analyzing the structure of attention in a transformer language model. In *Proceedings of the 2019 ACL Workshop BlackboxNLP: Analyzing and Interpreting Neural Networks for NLP*, pages 63–76.
- Vig, Jesse, Sebastian Gehrmann, Yonatan Belinkov, Sharon Qian, Daniel Nevo, Yaron Singer, and Stuart Shieber. 2020. Investigating gender bias in language models using causal mediation analysis. In *Advances in Neural Information Processing Systems*.
- Wahle, Jan Philip, Terry Ruas, Frederic Kirstein, and Bela Gipp. 2022. How large language models are transforming machine-paraphrase plagiarism. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 952–963.
- Wallace, Eric, Shi Feng, Nikhil Kandpal, Matt Gardner, and Sameer Singh. 2019a. Universal adversarial triggers for attacking and analyzing NLP. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 2153–2162.
- Wallace, Eric, Yizhong Wang, Sujian Li, Sameer Singh, and Matt Gardner. 2019b. Do NLP models know numbers? probing numeracy in embeddings. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 5307–5315.
- Wang, Alex, Yada Pruksachatkun, Nikita Nangia, Amanpreet Singh, Julian Michael, Felix Hill, Omer Levy, and Samuel Bowman. 2019. SuperGLUE: A stickier benchmark for general-purpose language understanding systems. In *Advances in Neural Information Processing Systems*, volume 32.
- Wang, Benyou, Lifeng Shang, Christina Lioma, Xin Jiang, Hao Yang, Qun Liu, and Jakob Grue Simonsen. 2021a. On position embeddings in BERT. In *International Conference on Learning Representations*.
- Wang, Boshi, Sewon Min, Xiang Deng, Jiaming Shen, You Wu, Luke Zettlemoyer, and Huan Sun. 2022a. Towards understanding chain-of-thought prompting: An empirical study of what matters. *ArXiv*, arXiv:2212.10001.
- Wang, Cunxiang, Boyuan Zheng, Yuchen Niu, and Yue Zhang. 2021b. Exploring generalization ability of pretrained language models on arithmetic and logical reasoning. In *Natural Language Processing and Chinese Computing*, pages 758–769, Springer International Publishing.
- Wang, Tianlu, Rohit Sridhar, Diyi Yang, and Xuezhi Wang. 2022b. Identifying and mitigating spurious correlations for improving robustness in NLP models. In *Findings of the Association for Computational Linguistics: NAACL 2022*, pages 1719–1729.
- Warstadt, Alex, Yu Cao, Ioana Grosu, Wei Peng, Hagen Blix, Yining Nie, Anna Alsop, Shikha Bordia, Haokun Liu, Alicia Parrish, Sheng-Fu Wang, Jason Phang, Anhad Mohananey, Phu Mon Htut, Paloma Jeretic, and Samuel R. Bowman. 2019. Investigating BERT's knowledge of language: Five analysis methods with NPIs. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 2877–2887.
- Warstadt, Alex, Alicia Parrish, Haokun Liu, Anhad Mohananey, Wei Peng, Sheng-Fu Wang, and Samuel R. Bowman. 2020. BLiMP: The benchmark of linguistic minimal pairs for English. *Transactions of the Association for Computational Linguistics*, 8:377–392.
- Webb, Taylor, Keith Holyoak, and Hongjing Lu. 2022. Emergent analogical reasoning in large language models. *ArXiv*, arXiv:2212.09196.
- Wei, Jason, Maarten Bosma, Vincent Zhao, Kelvin Guu, Adams Wei Yu, Brian Lester,

- Nan Du, Andrew M. Dai, and Quoc V. Le. 2022a. Finetuned language models are zero-shot learners. In *International Conference on Learning Representations*.
- Wei, Jason, Dan Garrette, Tal Linzen, and Ellie Pavlick. 2021. Frequency effects on syntactic rule learning in transformers. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 932–948.
- Wei, Jason, Najoung Kim, Yi Tay, and Quoc V. Le. 2023. Inverse scaling can become U-shaped. *ArXiv*, arXiv:2211.02011.
- Wei, Jason, Yi Tay, Rishi Bommasani, Colin Raffel, Barret Zoph, Sebastian Borgeaud, Dani Yogatama, Maarten Bosma, Denny Zhou, Donald Metzler, Ed H. Chi, Tatsunori Hashimoto, Oriol Vinyals, Percy Liang, Jeff Dean, and William Fedus. 2022b. Emergent abilities of large language models. *Transactions on Machine Learning Research*.
- Wei, Jason, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed Chi, Quoc Le, and Denny Zhou. 2022c. Chain of thought prompting elicits reasoning in large language models. In *Advances in Neural Information Processing Systems*.
- Weidinger, Laura, John F. J. Mellor, Maribeth Rauh, Conor Griffin, Jonathan Uesato, Po-Sen Huang, Myra Cheng, Mia Glaese, Borja Balle, Atoosa Kasirzadeh, Zachary Kenton, Sande Minnich Brown, William T. Hawkins, Tom Stepleton, Courtney Biles, Abeba Birhane, Julia Haas, Laura Rimell, Lisa Anne Hendricks, William S. Isaac, Sean Legassick, Geoffrey Irving, and Iason Gabriel. 2021. Ethical and social risks of harm from language models. *ArXiv*, arXiv:2112.04359.
- Weidinger, Laura, Jonathan Uesato, Maribeth Rauh, Conor Griffin, Po-Sen Huang, John Mellor, Amelia Glaese, Myra Cheng, Borja Balle, Atoosa Kasirzadeh, Courtney Biles, Sasha Brown, Zac Kenton, Will Hawkins, Tom Stepleton, Abeba Birhane, Lisa Anne Hendricks, Laura Rimell, William Isaac, Julia Haas, Sean Legassick, Geoffrey Irving, and Iason Gabriel. 2022. Taxonomy of risks posed by language models. In *Proceedings of the ACM Conference on Fairness, Accountability, and Transparency*, page 214–229, Association for Computing Machinery, New York, NY, USA.
- Weir, Nathaniel, Adam Poliak, and Benjamin Van Durme. 2020. Probing neural language models for human tacit assumptions. In *Annual Meeting of the Cognitive Science Society*, volume 42.
- Weissweiler, Leonie, Valentin Hofmann, Abdullatif Köksal, and Hinrich Schütze. 2022. The better your syntax, the better your semantics? probing pretrained language models for the English comparative correlative. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 10859–10882.
- Wettig, Alexander, Tianyu Gao, Zexuan Zhong, and Danqi Chen. 2023. Should you mask 15% in masked language modeling? In *Proceedings of the Conference of the European Chapter of the Association for Computational Linguistics*.
- White, Jennifer C. and Ryan Cotterell. 2021. Examining the inductive bias of neural language models with artificial languages. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 454–463.
- Wilcox, Ethan, Ethan Gotlieb, Richard Futrell, and Roger Levy. 2022. Using Computational Models to Test Syntactic Learnability. *Linguistic Inquiry*, pages 1–88.
- Williams, Adina, Nikita Nangia, and Samuel Bowman. 2018. A broad-coverage challenge corpus for sentence understanding through inference. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers)*, pages 1112–1122.
- Wu, Shijie and Mark Dredze. 2020. Are all languages created equal in multilingual BERT? In *Proceedings of the 5th Workshop on Representation Learning for NLP*, pages 120–130.
- Xia, Mengzhou, Mikel Artetxe, Chunting Zhou, Xi Victoria Lin, Ramakanth Pasunuru, Danqi Chen, Luke Zettlemoyer, and Ves Stoyanov. 2022. Training trajectories of language models across scales. *ArXiv*, arXiv:2212.09803.
- Yu, Charles, Ryan Sie, Nicolas Tedeschi, and Leon Bergen. 2020. Word frequency does not predict grammatical knowledge in language models. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 4040–4054.
- Zafriir, Ofir, Guy Boudoukh, Peter Izsak, and Moshe Wasserblat. 2019. Q8BERT: Quantized 8bit BERT. In *Fifth Workshop on Energy Efficient Machine Learning and*

- Cognitive Computing*, pages 36–39, IEEE.
- Zellers, Rowan, Ari Holtzman, Elizabeth Clark, Lianhui Qin, Ali Farhadi, and Yejin Choi. 2021. TuringAdvice: A generative and dynamic evaluation of language use. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 4856–4880.
- Zhang, Li, Hai Xu, Yue Yang, Shuyan Zhou, Weiqiu You, Manni Arora, and Chris Callison-Burch. 2023. Causal reasoning of entities and events in procedural texts. In *Findings of the Association for Computational Linguistics: EACL*.
- Zhang, Lining, Mengchen Wang, Liben Chen, and Wenxin Zhang. 2022a. Probing GPT-3’s linguistic knowledge on semantic tasks. In *Proceedings of the Fifth BlackboxNLP Workshop on Analyzing and Interpreting Neural Networks for NLP*, pages 297–304.
- Zhang, Sheng, Xin Zhang, Weiming Zhang, and Anders Søgaard. 2021a. Sociolectal analysis of pretrained language models. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 4581–4588.
- Zhang, Susan, Stephen Roller, Naman Goyal, Mikel Artetxe, Moya Chen, Shuohui Chen, Christopher Dewan, Mona Diab, Xian Li, Xi Victoria Lin, Todor Mihaylov, Myle Ott, Sam Shleifer, Kurt Shuster, Daniel Šimig, Punit Singh Koura, Anjali Sridhar, Tianlu Wang, and Luke Zettlemoyer. 2022b. OPT: Open pre-trained Transformer language models. *ArXiv*, arXiv:2205.01068.
- Zhang, Tianyi, Varsha Kishore, Felix Wu, Kilian Q. Weinberger, and Yoav Artzi. 2020. BERTScore: Evaluating text generation with BERT. In *International Conference on Learning Representations*.
- Zhang, Yian, Alex Warstadt, Xiaocheng Li, and Samuel R. Bowman. 2021b. When do you need billions of words of pretraining data? In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 1112–1125.
- Zhang, Zhengyan, Xu Han, Zhiyuan Liu, Xin Jiang, Maosong Sun, and Qun Liu. 2019. ERNIE: Enhanced language representation with informative entities. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 1441–1451.
- Zhao, Yiyun, Jian Gang Ngui, Lucy Hall Hartley, and Steven Bethard. 2021. Do pretrained transformers infer telicity like humans? In *Proceedings of the 25th Conference on Computational Natural Language Learning*, pages 72–81.
- Zhao, Zhixue, Ziqi Zhang, and Frank Hopfgartner. 2021. A comparative study of using pre-trained language models for toxic comment classification. In *The ACM Web Conference*, page 500–507, Association for Computing Machinery, New York, NY, USA.
- Zhou, Kaitlyn, Kawin Ethayarajh, and Dan Jurafsky. 2022. Richer countries and richer representations. In *Findings of the Association for Computational Linguistics: ACL 2022*, pages 2074–2085.
- Zhou, Pei, Rahul Khanna, Seyeon Lee, Bill Yuchen Lin, Daniel Ho, Jay Pujara, and Xiang Ren. 2021. RICA: Evaluating robust inference capabilities based on commonsense axioms. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 7560–7579.
- Zong, Mingyu and Bhaskar Krishnamachari. 2022. A survey on GPT-3. *ArXiv*, arXiv:2212.00857.
- Zwaan, Rolf A. 2016. Situation models, mental simulations, and abstract concepts in discourse comprehension. *Psychonomic Bulletin & Review*, 23:1028–1034.

Appendix A: Literature Review Process

We identified papers to include in this survey using Semantic Scholar (Fricke 2018). From a seed of 271 relevant language model analysis papers (including the majority of the citation list from Rogers, Kovaleva, and Rumshisky 2020), we extracted all papers that cited any paper in the seed. This resulted in over 15K papers, last scraped on February 4, 2023. Anecdotally, the majority of recent language model analysis papers we encountered were included in this list. We manually filtered by title down to approximately 1500 potentially relevant papers, gradually refining the scope as described in Section 1.1. We then further filtered by abstract down to approximately 400 highly relevant papers.